

INTRODUCTION TO DEEP LEARNING

UPC TelecomBCN Barcelona (4th edition). Spring Edition.



Instructors:



Xavier
Giró-i-Nieto



Ferran
Marqués



Ramon
Morros



Montse
Pardàs



Javier
Ruiz



Elisa
Sayrol



Verònica
Vilaplana

Teaching Assistants:



Gerard
Gallego



Albert
Mosella

Day 6 Lecture 2

Recurrent Neural Networks



Xavier Giro-i-Nieto



@DocXavi



xavier.giro@upc.edu



Associate Professor
Universitat Politècnica de Catalunya

Acknowledgments



Santiago Pascual

A screenshot of a website for a winter seminar. The header reads "Winter Seminar UPC TelecomBCN 24 - 25 January 2017". Below it, "Day 2 Lecture 2" and "Recurrent Neural Networks I" are prominently displayed. A video player shows a video of a speaker. On the right, there is a profile picture of Santiago Pascual and some social media icons.



Marta R. Costa-jussà

A screenshot of the same website as the previous slide, but for Day 2 Lecture 3. The title "Recurrent Neural Networks II" is shown. A video player shows a video of a speaker. On the right, there is a profile picture of Santiago Pascual and some social media icons.

A presentation slide with a light blue header containing the title "Recurrent Neural Networks" and the name "DLAI - MARTA R. COSTA-JUSSÀ". The main content area is white. In the bottom right corner, there is a small video frame showing a person standing in front of a whiteboard, and the footer includes the UPC logo and text about the Department of Teoria del Senyal i Comunicacions.



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH
Language and Speech Technologies



Santiago Pascual
UPC TelecomBCN
DLSL 2017

Marta R. Costa-Jussà
UPC TelecomBCN
DLAI 2017

Similar Lecture



Lecture 11

Recurrent Neural Networks

DEEP LEARNING FOR ARTIFICIAL INTELLIGENCE
Masters @ UPC TelecomBCN Barcelona (4th edition), Autumn 2020.

Instructors:

- Xavier Giro-i-Nieto
- Ventura Vilaplana
- Javier Perea
- Ramon Martínez
- Albert Moselha
- Luis Selgasero

Guest:

- Yannick Katsoulis (River Labs)

[course site]



Xavier Giro-i-Nieto

Associate Professor
Universitat Politècnica de Catalunya

[@DocXavi](#)
 xavier.giro@upc.edu



**What's your
favourite series ?**



Time Series



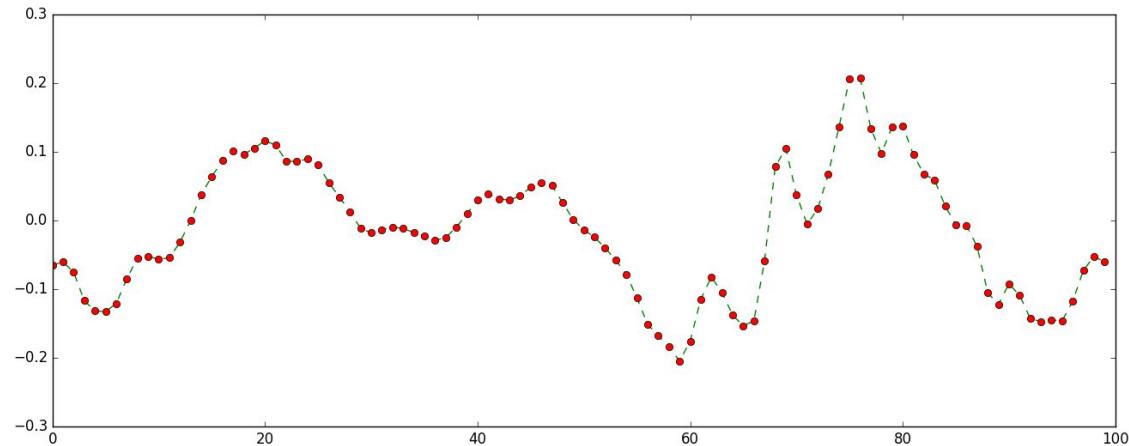
Source: ([AI Memes](#)) AI & Deep Learning Memes For Back-propagated Poets

Outline

1. Motivation

Sequences (a.k.a time series)

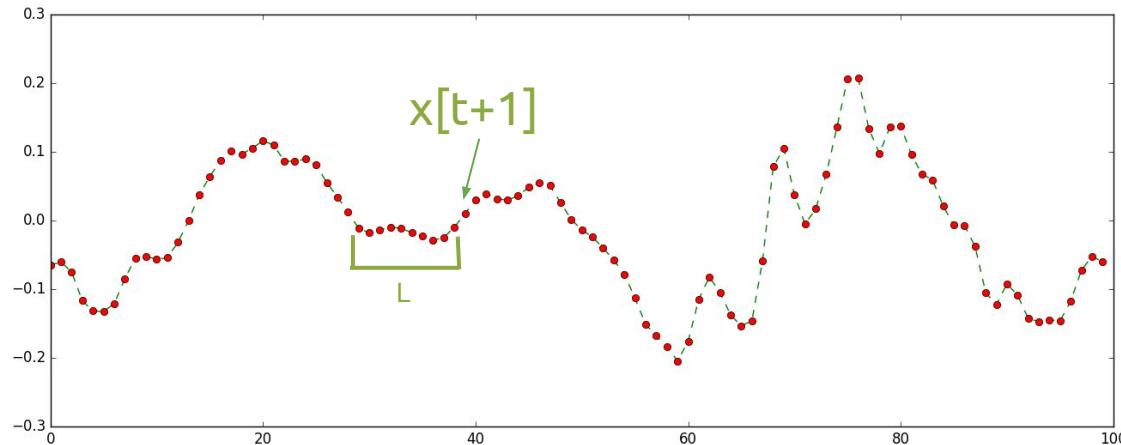
Consider a sequence of samples (eg. of temperature):



Sequences with Feed Forward MLP

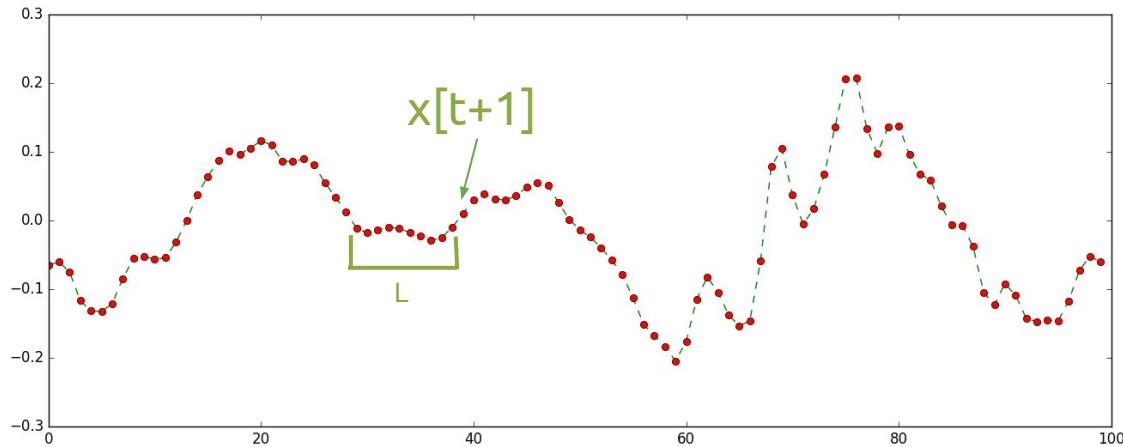
Task:

Predict sample $x[t+1]$ knowing previous values $\{x[t], x[t-1], x[t-2], \dots, x[t-L]\}$



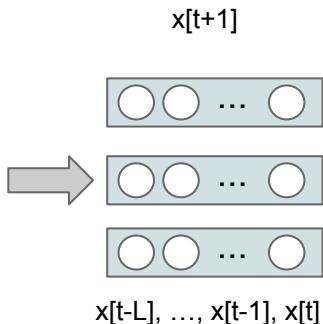
How could you solve this task with a naive feed forward MLP ?

Sequences with Feed Forward MLP

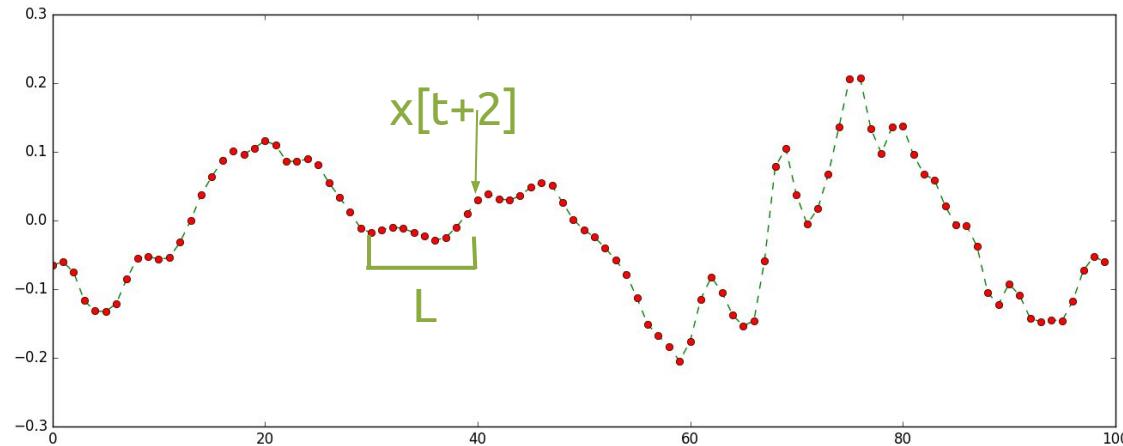


Feed Forward approach:

- static window of size L
- slide the window time-step wise

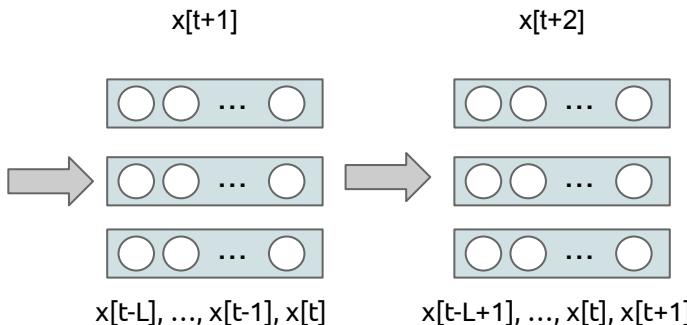


Sequences with Feed Forward MLP

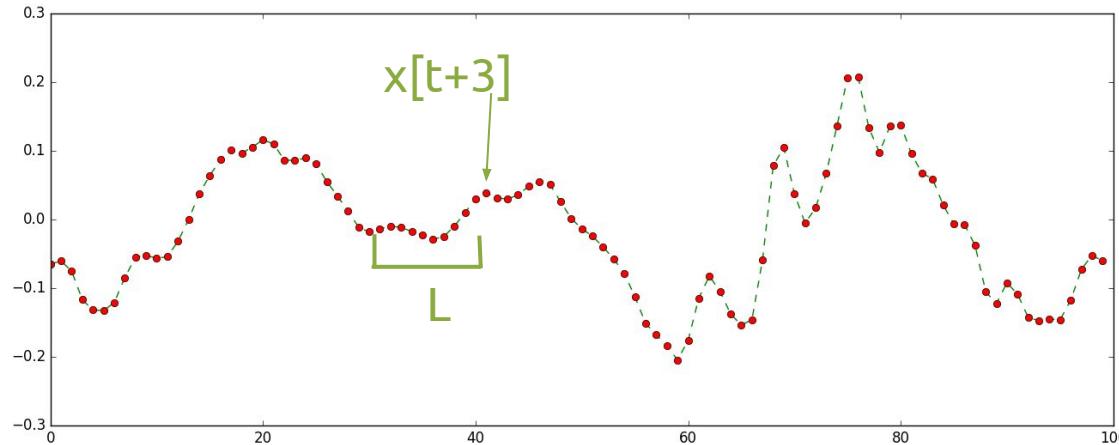


Feed Forward approach:

- static window of size L
- slide the window time-step wise

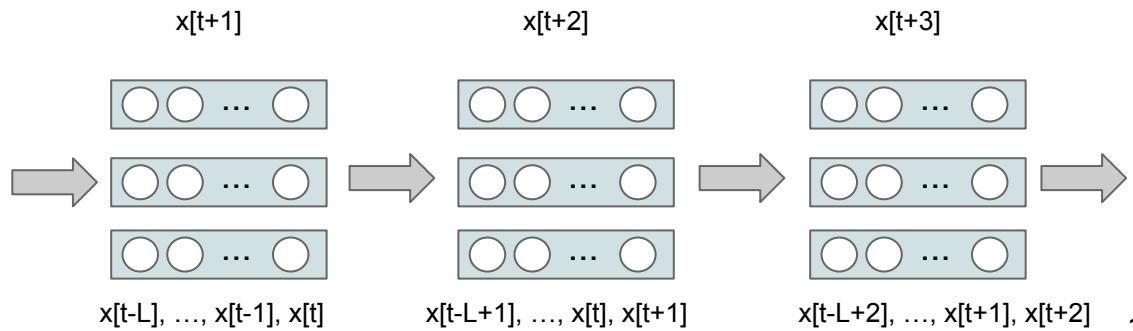


Sequences with Feed Forward MLP



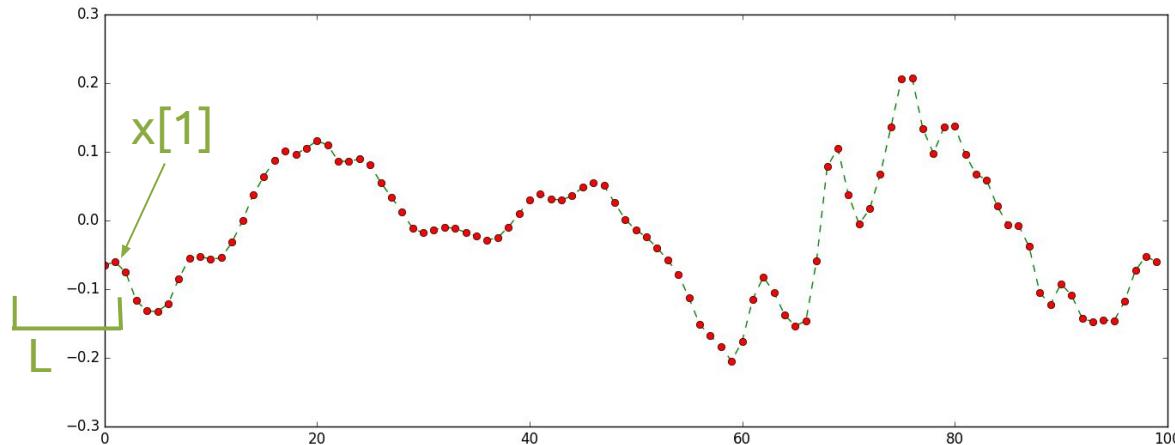
Feed Forward approach:

- static window of size L
- slide the window time-step wise



Sequences with Feed Forward MLP

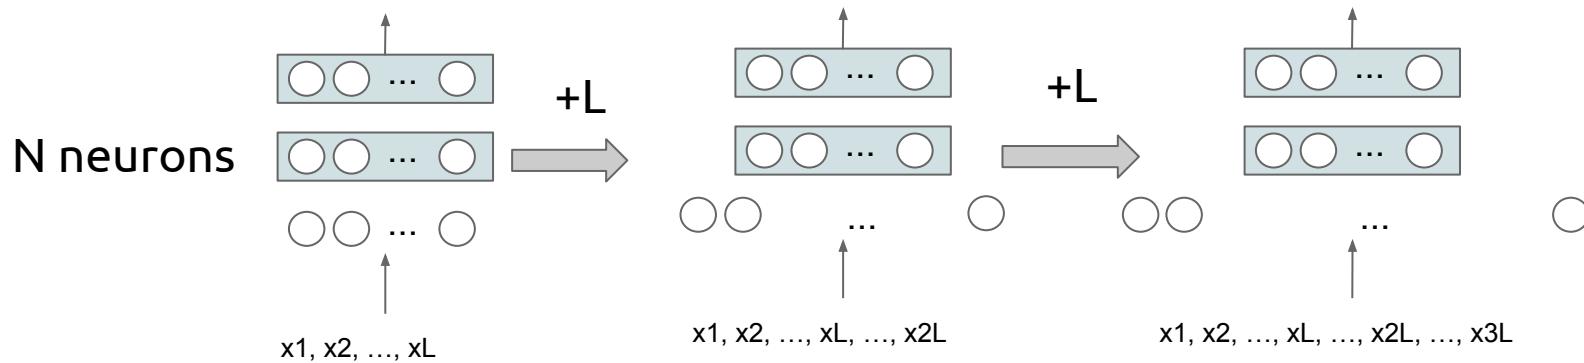
How would you deal with the prediction at the beginning of the sequence ?



By adding (zero) padding for the samples necessary to fill the input layer of the MLP.

Sequences with Feed Forward MLP

In the first layer, how does an increase of the receptive field (L) affect the amount of parameters to learn ?



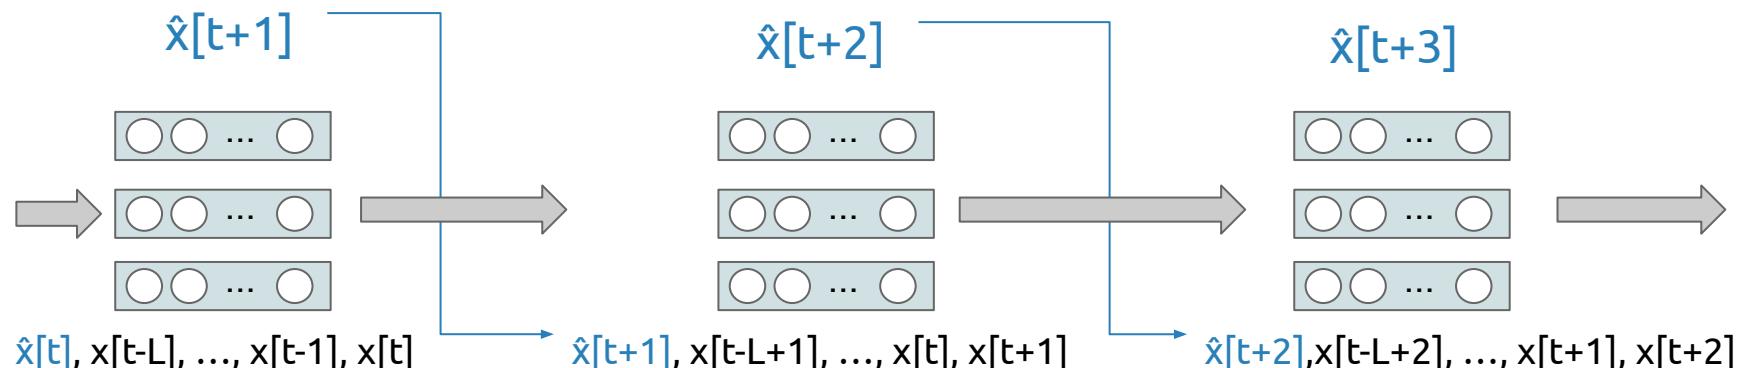
Receptive field	L	$2L$	$3L$
#params in 1st layer	$N \times L + N$	$N \times 2L + N$	$N \times 3L + N$

A $+1$ increase of the receptive field increases the parameters with a factor N .

Sequences with Feed Forward MLP

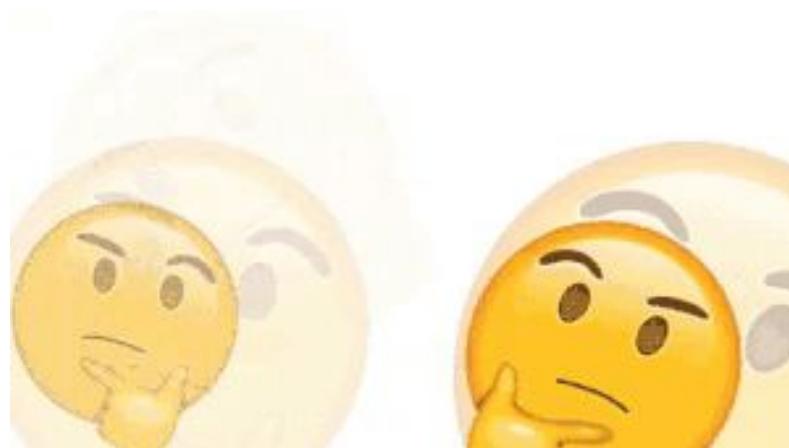
How could the previous decisions condition the future predictions ?

Adopting an **auto-regressive** strategy, by which the previous outputs become inputs in future steps.



Could we find an architecture that:

- adapts to any input length to avoid padding.
- does not increase the complexity with the receptive field.
- can naturally remember what it has seen in previous timesteps ?



Outline

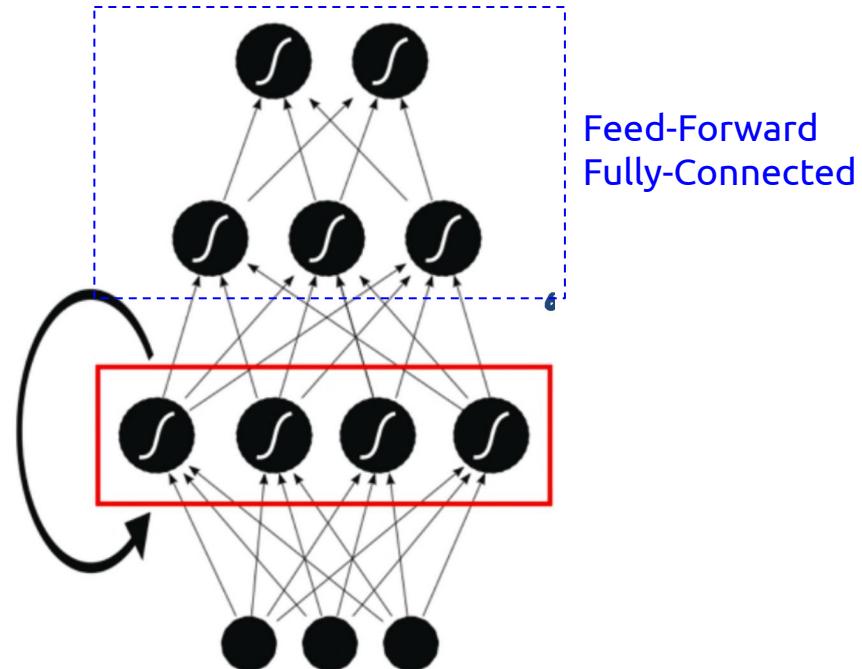
1. Motivation
2. Recurrent Neural Layer

Recurrent Neural Layer

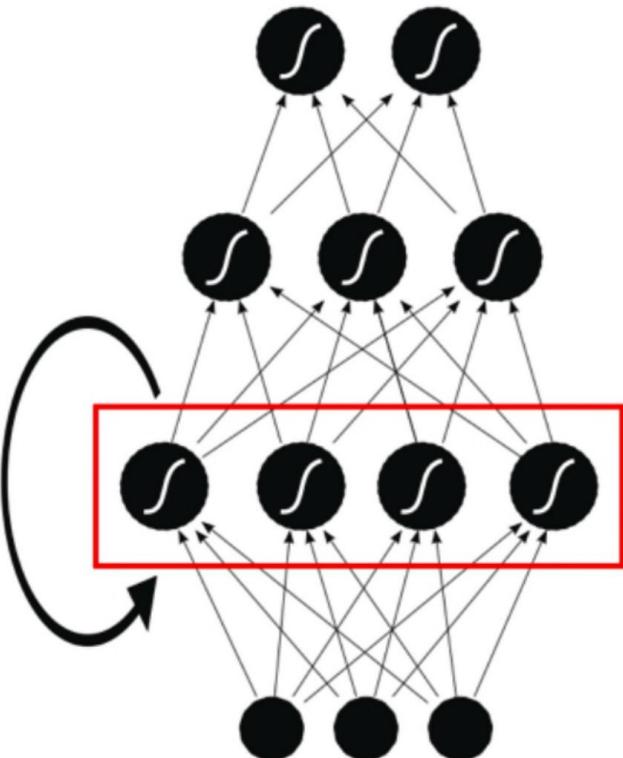


Recurrent layer (RNN)

The hidden state of a recurrent layer h_t also depends from its previous state h_{t-1} .



Recurrent Neural Layer



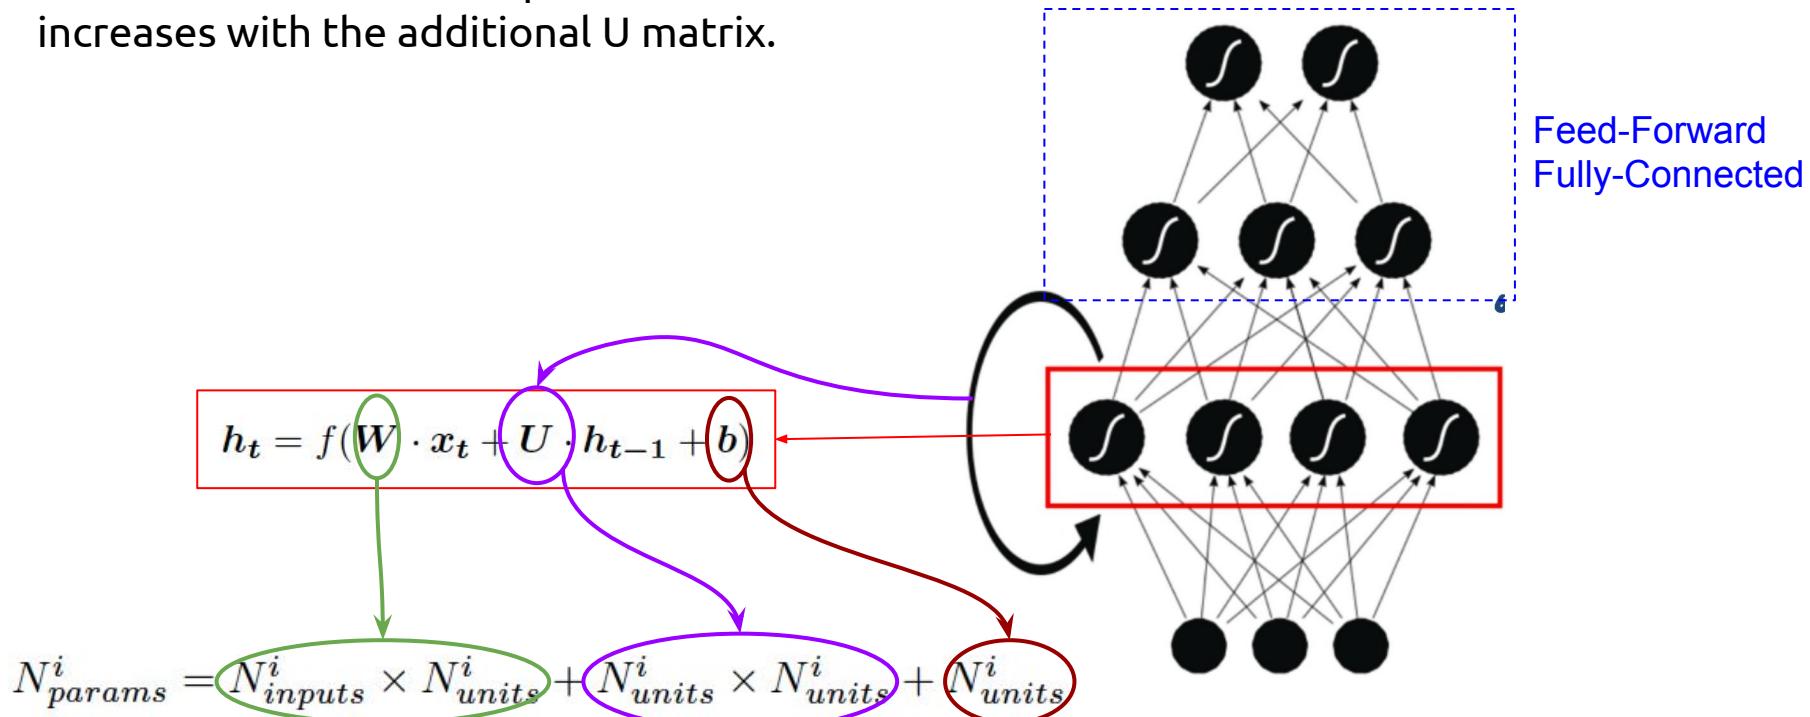
Feed-forward
Weights (W)

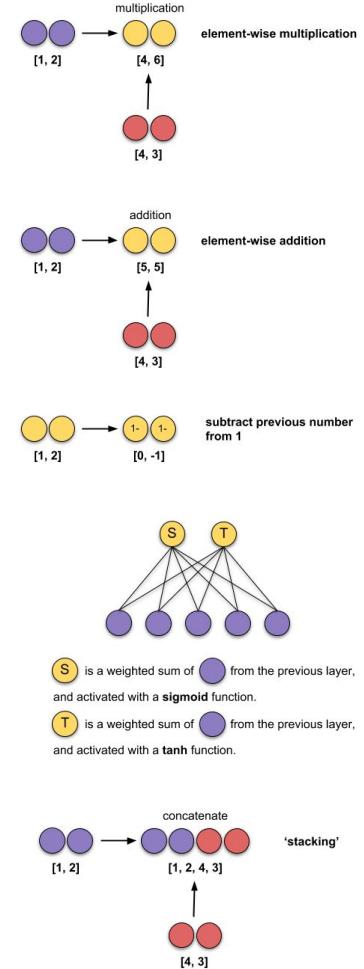
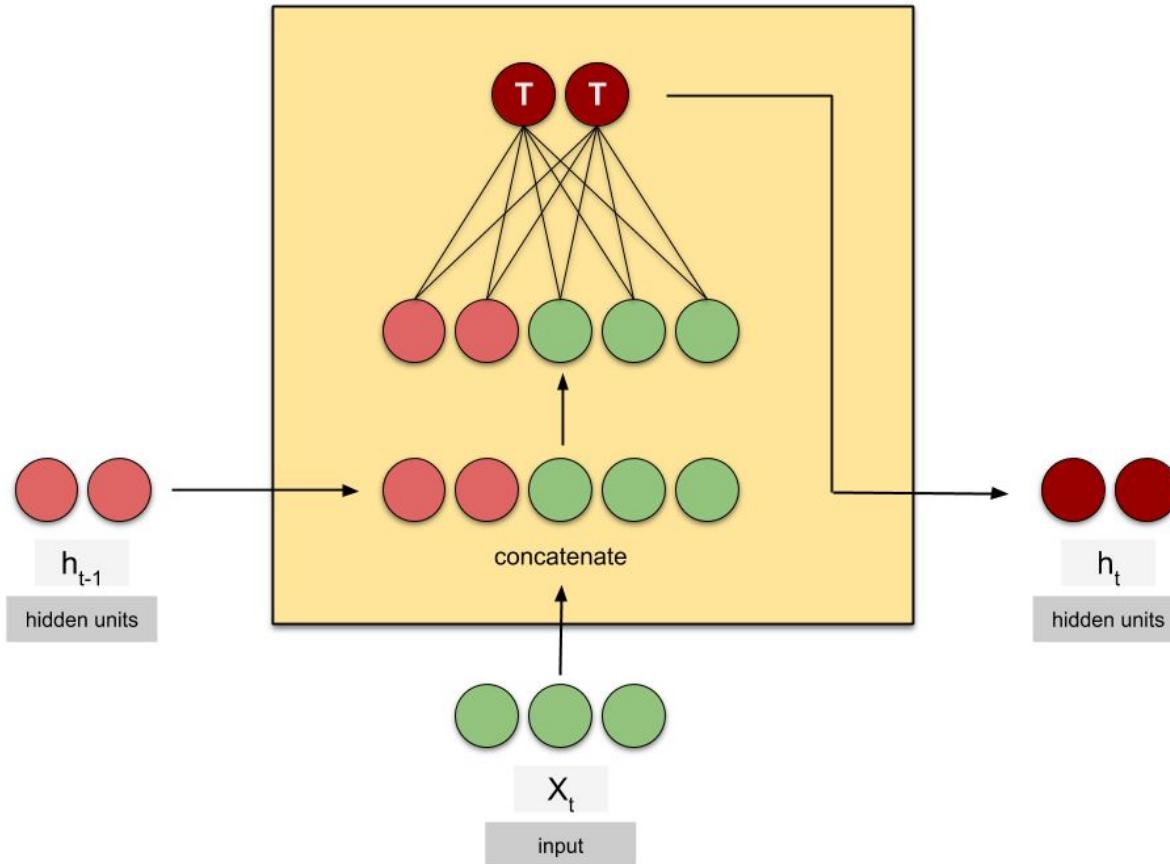
$$h_t = f(W \cdot x_t + U \cdot h_{t-1} + b)$$

Recurrent
Weights (U)

Recurrent Neural Layer

The amount of learnable parameters increases with the additional U matrix.



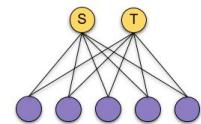
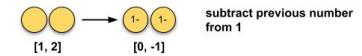
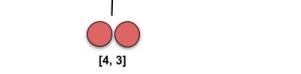
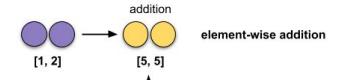
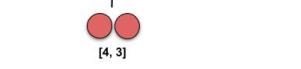
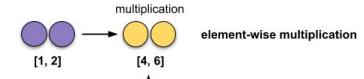
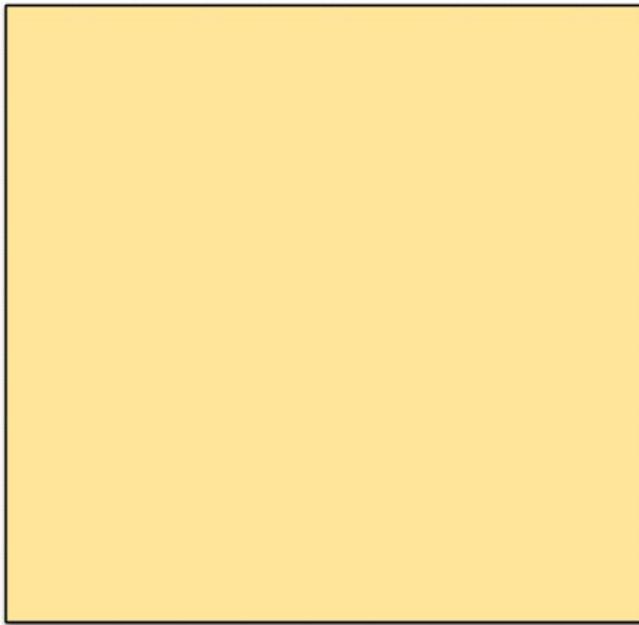




h_{t-1}
hidden units

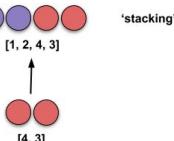
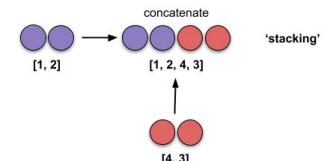


X_t
input



S is a weighted sum of h_{t-1} and activated with a **sigmoid** function.

T is a weighted sum of h_{t-1} and activated with a **tanh** function.

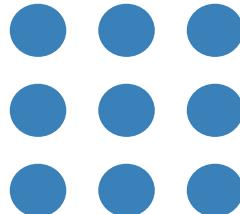


Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time

Forward Pass in RNN

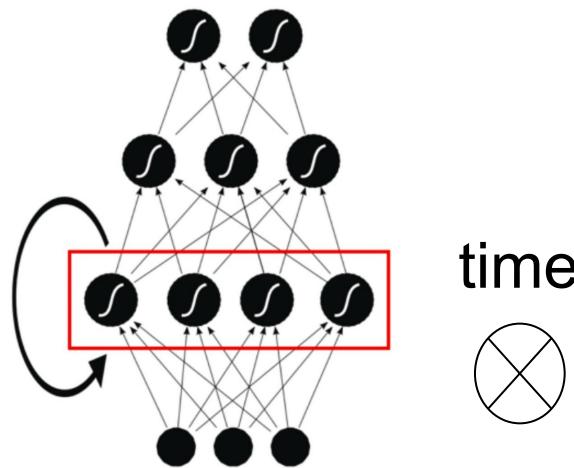
Front View



Side View

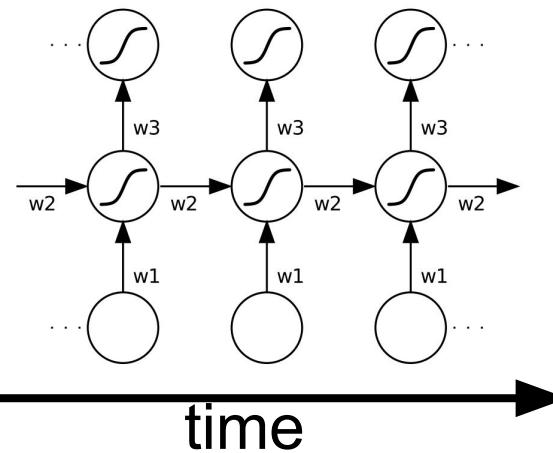


Rotation
90°



time

Unfold
(Rotation
90°)

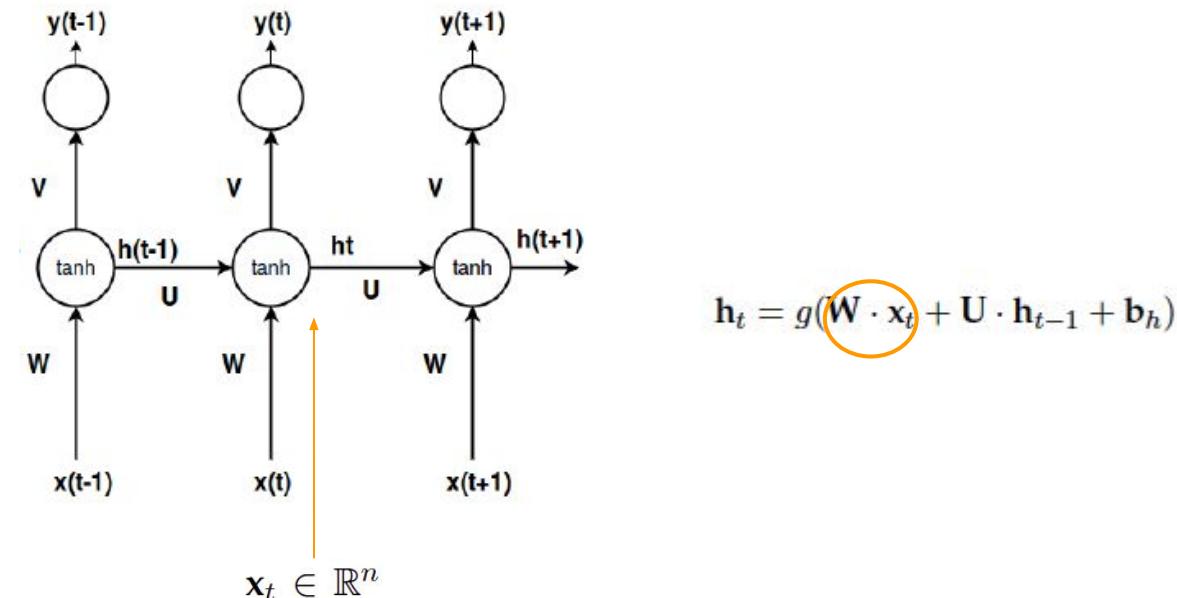


time

Forward Pass in RNN

Hence we have two data flows: **Forward in neural layers + time** propagation

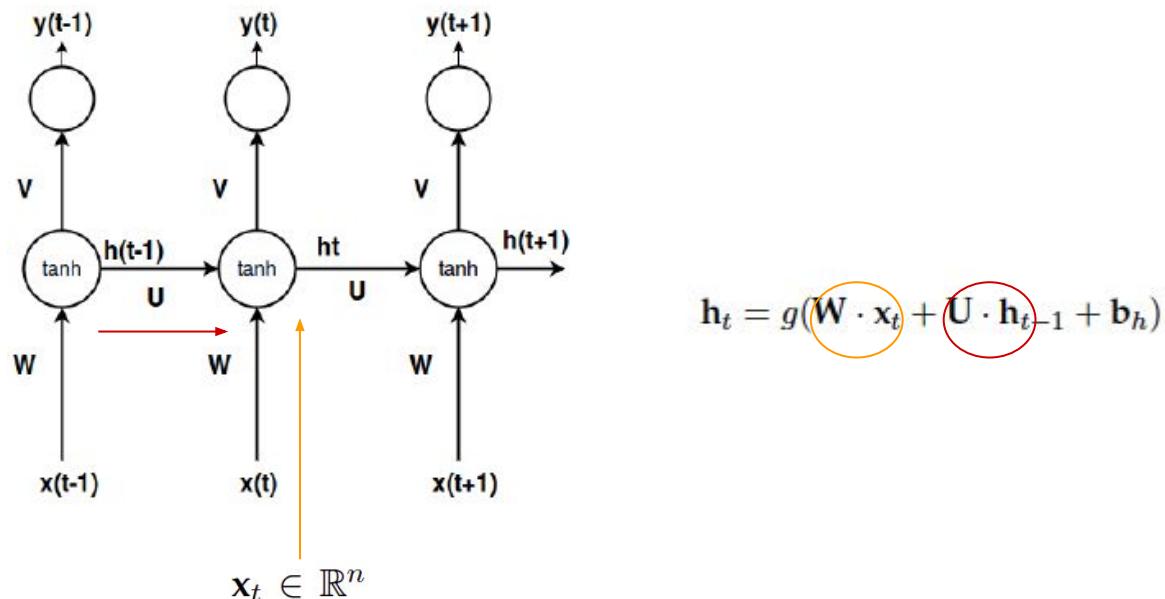
BEWARE: We have extra temporal depth now! Every time-step is an extra level of depth (as a deeper stack of layers in a feed-forward fashion!)



Forward Pass in RNN

Hence we have two data flows: **Forward in layers + time** propagation

- Last time-step includes the context of our decisions recursively

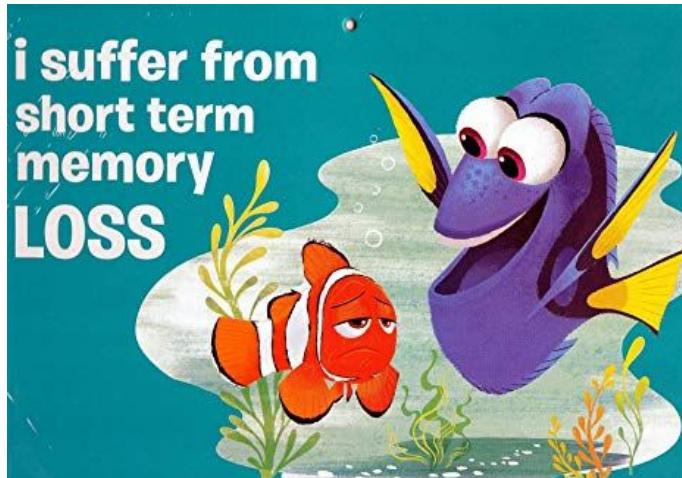


Forward Pass in RNN

Main problem:

- **Long-term memory** (remembering quite far time-steps) **vanishes quickly** because of the recursive operation with non-linearities $\mathbf{g}(\cdot)$ and \mathbf{U} .

$$\mathbf{h}_t = \mathbf{g}(\mathbf{W} \cdot \mathbf{x}_t + \mathbf{U} \cdot \mathbf{g}(\cdots \mathbf{g}(\mathbf{W} \cdot \mathbf{x}_{t-T} + \mathbf{U} \cdot \mathbf{h}_{t-T} + \mathbf{b}_h) \cdots) + \mathbf{b}_h)$$



Backpropagation Through Time (BPTT)

Back Propagation Through Time (BPTT): The training method has to take into account the time operations:

Total error at the output is the sum of errors at each time-step t

$$E(\mathbf{y}, \hat{\mathbf{y}}) = \sum_{t=1}^T E_t(\mathbf{y}_t, \hat{\mathbf{y}}_t)$$

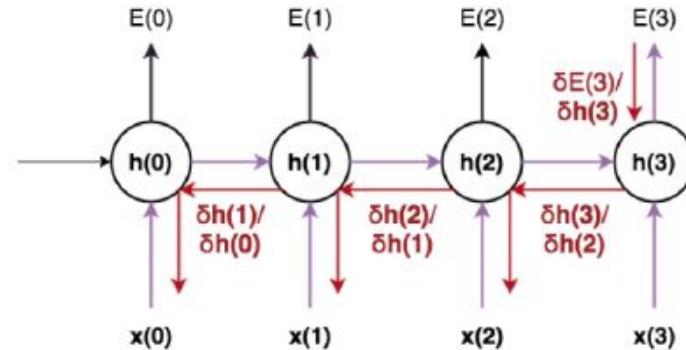
Total gradient is the sum of gradients at each time-step t

$$\frac{\partial E}{\partial \mathbf{W}} = \sum_{t=0}^{T-1} \frac{\partial E_t}{\partial \mathbf{W}}$$

T: max amount of time-steps to do back-prop. In Keras this is specified when defining the “input shape” to the RNN layer, by means of:
(batch size, sequence length (T), input_dim)

Input shape

3D tensor with shape `(nb_samples, timesteps, input_dim)`.



Example back-prop in time with 3 time-steps

Backpropagation Through Time (BPTT)

Main problems:

- **Exploding / vanishing gradients:** During training gradients explode/vanish easily because of depth-in-time.

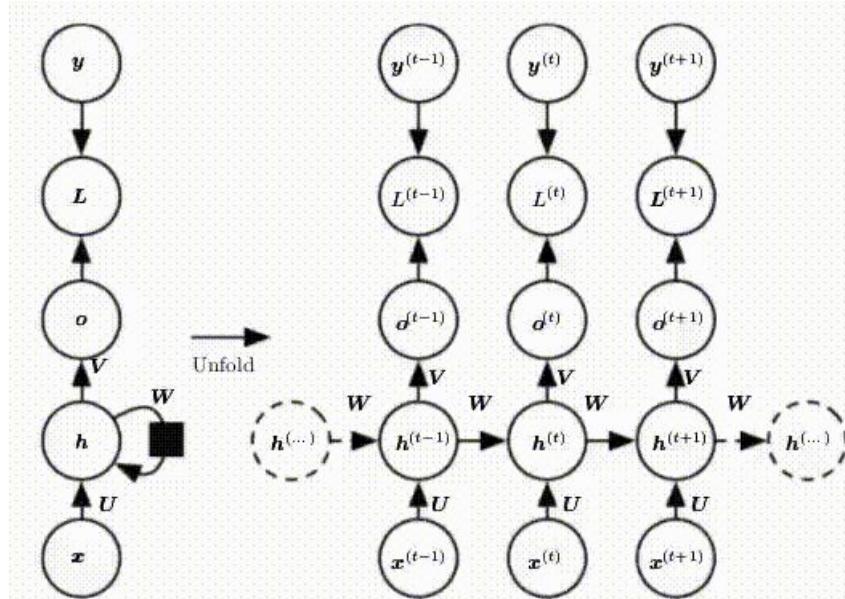


Figure: [Jordi Pons](#)

Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time
4. Gated Recurrent Neurons
 - a. LSTM
 - b. GRU

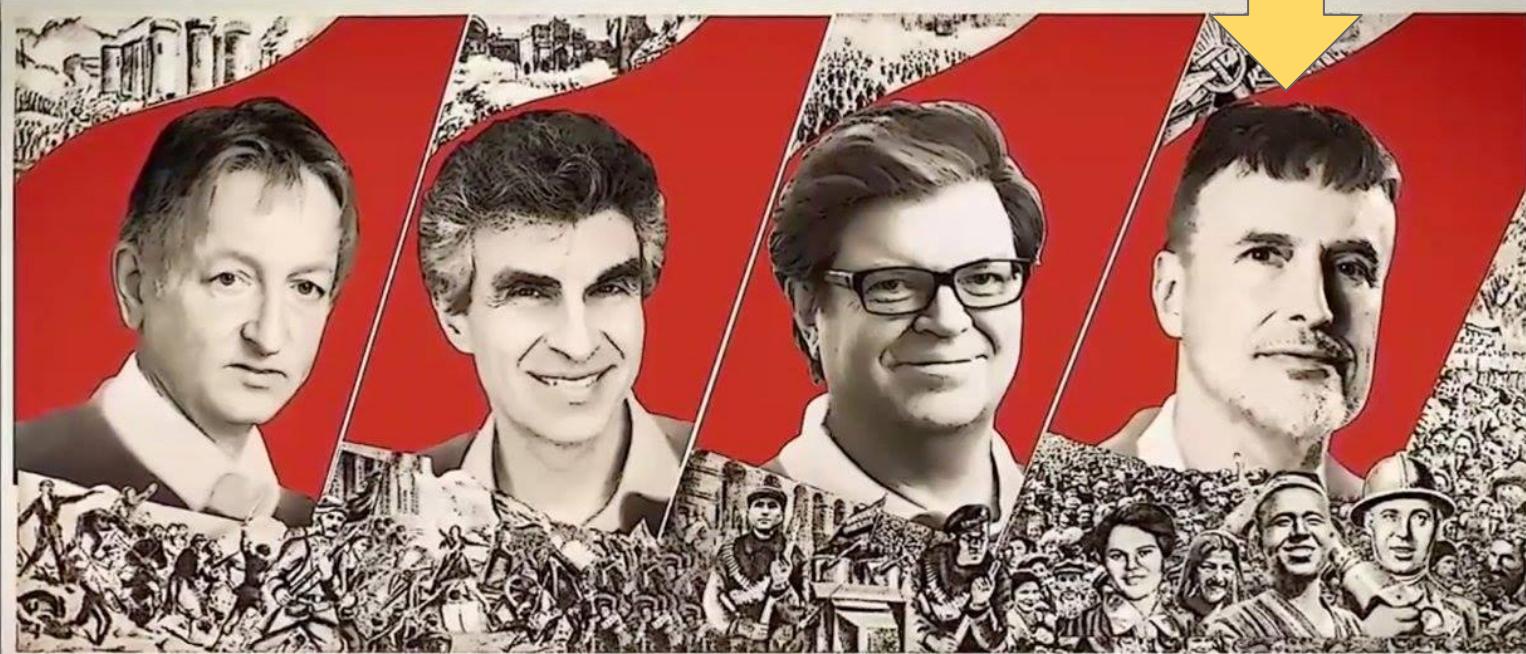


Figure: Michael Bronstein, "[Geometric Deep Learning](#)" (ICLR 2021)

Long Short-Term Memory (LSTM)

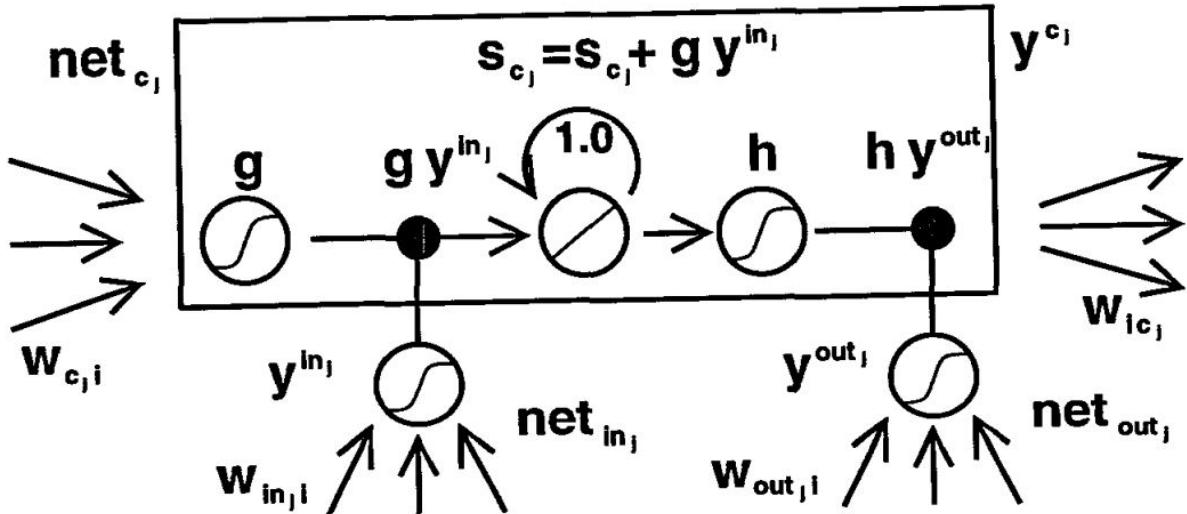
The New York Times, ["When A.I. Matures, It May Call Jürgen Schmidhuber 'Dad'"](#)
(November 2016)



Long Short-Term Memory (LSTM)

1744

Sepp Hochreiter and Jürgen Schmidhuber



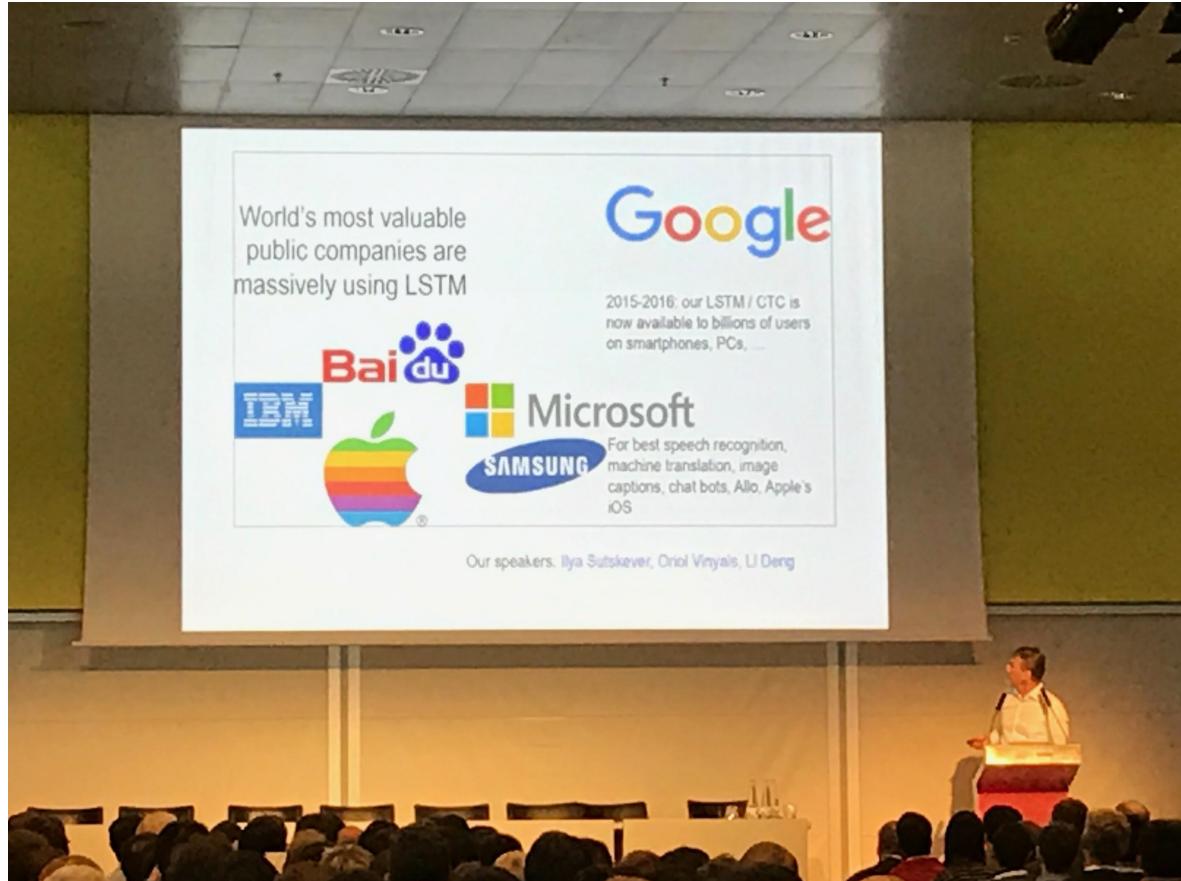
Hochreiter, Sepp, and Jürgen Schmidhuber. ["Long short-term memory."](#) Neural computation 9, no. 8 (1997): 1735-1780.

Long Short-Term Memory (LSTM)

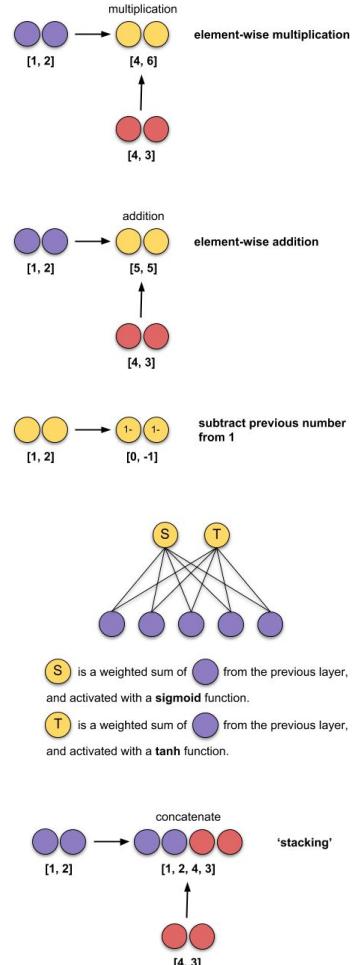
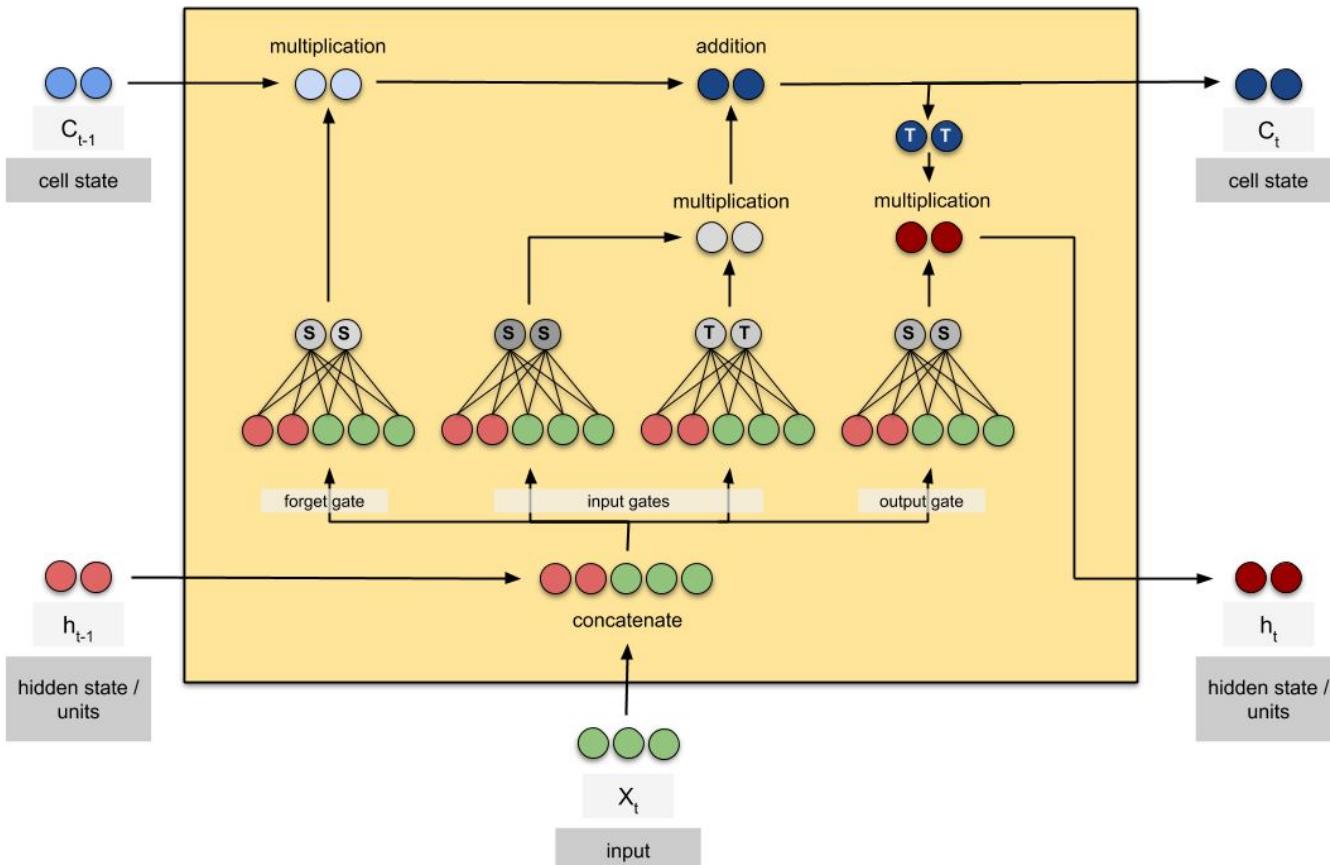


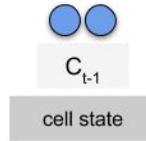
Jürgen
Schmidhuber @
NIPS 2016
Barcelona

Long Short-Term Memory (LSTM)



Jürgen Schmidhuber
@ NIPS 2016 Barcelona





C_{t-1}

cell state



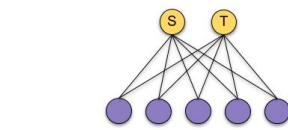
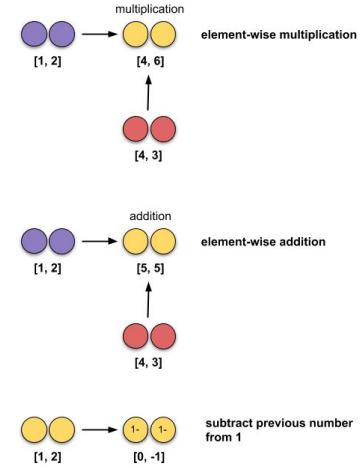
h_{t-1}

hidden state / units



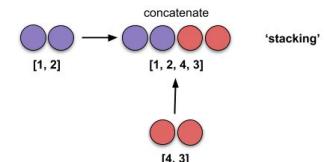
X_t

input



S is a weighted sum of h_{t-1} from the previous layer, and activated with a **sigmoid** function.

T is a weighted sum of h_{t-1} from the previous layer, and activated with a **tanh** function.



Gating method



Solutions:

1. Change the way in which past information is kept → create the notion of **cell state**: a memory unit that keeps long-term information in a safer way by protecting it from recursive operations.
2. Make every RNN unit able to **forget whatever may not be useful anymore** by clearing that info from the cell state (optimized clearing mechanism)
3. Make every RNN unit able to decide whether **the current time-step information matters or not**, to accept or discard (optimized reading mechanism)
4. Make every RNN unit able to **output the decisions whenever it is ready to do so** (optimized output mechanism)

Long Short-Term Memory (LSTM)

Three **gates** are governed by *sigmoid* units (btw [0,1]) define the control of in & out information with a product..

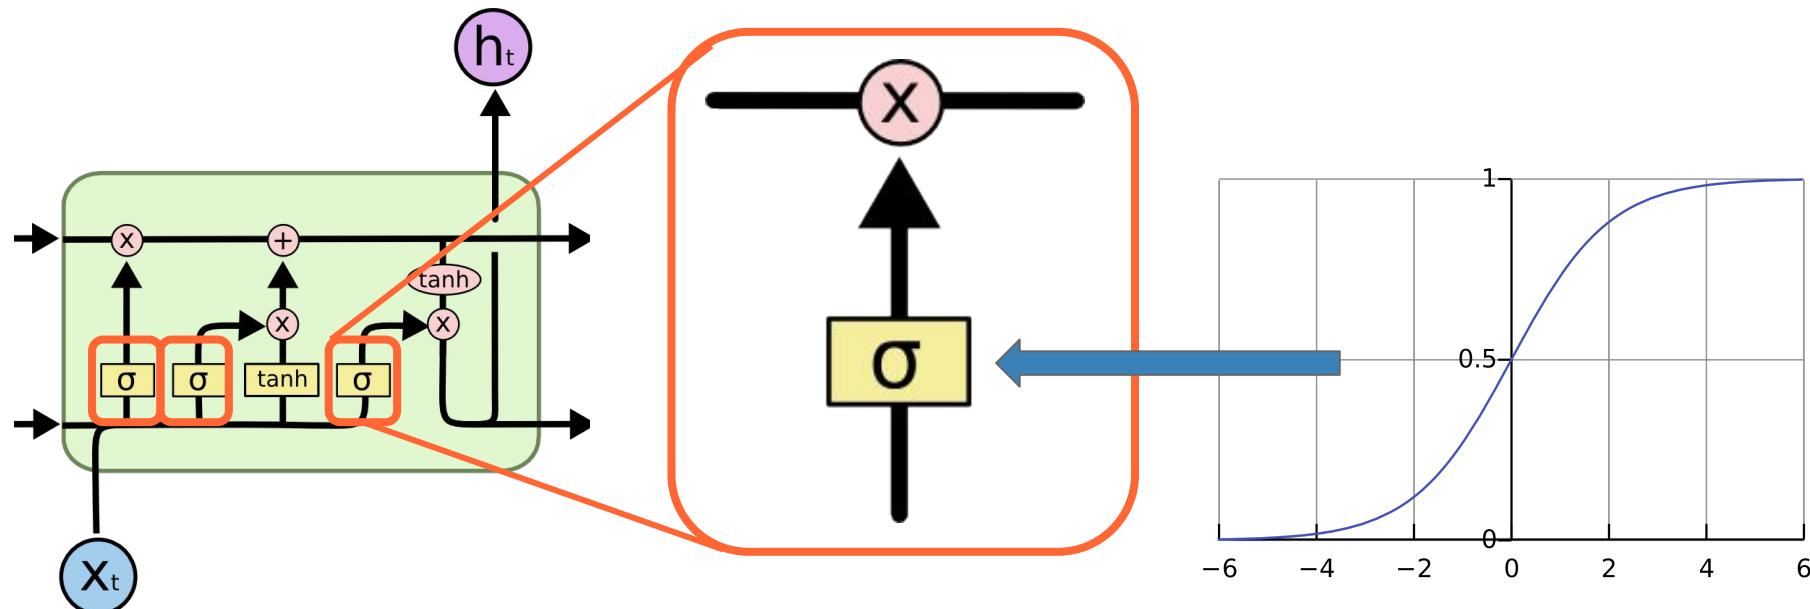
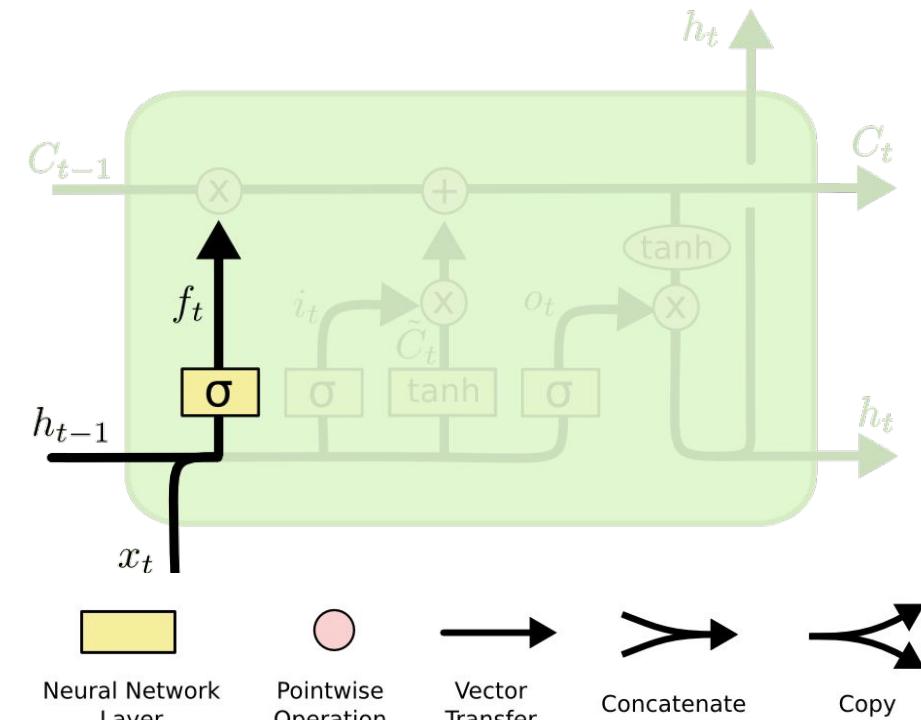


Figure: Cristopher Olah, ["Understanding LSTM Networks"](#) (2015)

Long Short-Term Memory (LSTM)

Make every RNN unit able to **forget whatever may not be useful anymore** by clearing that info from the cell state (optimized clearing mechanism)



Forget Gate:

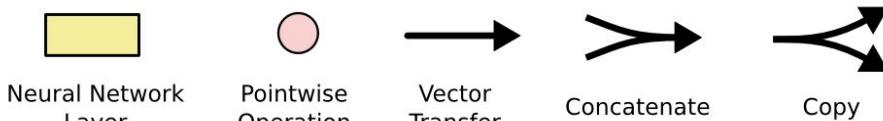
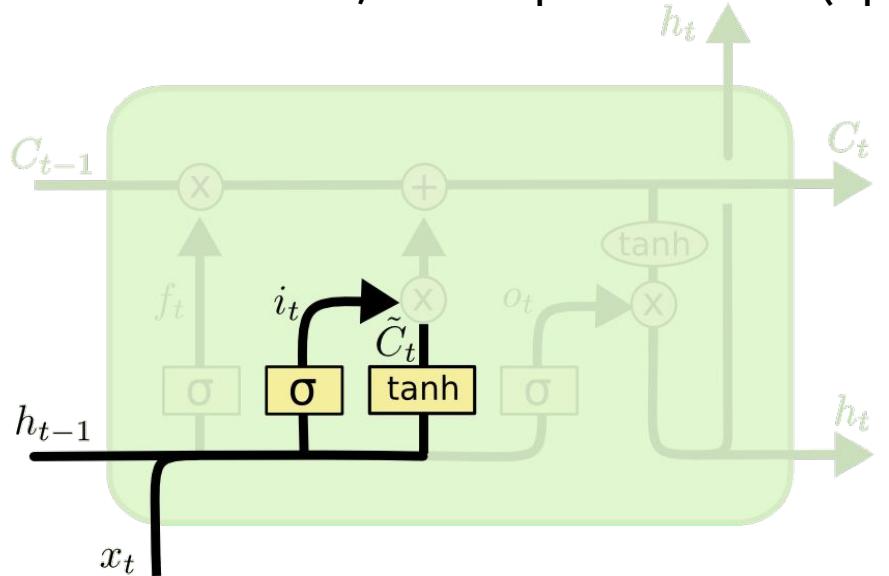
$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$



Concatenate

Long Short-Term Memory (LSTM)

Make every RNN unit able to decide whether **the current time-step information matters or not**, to accept or discard (optimized reading mechanism)



Input Gate Layer

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

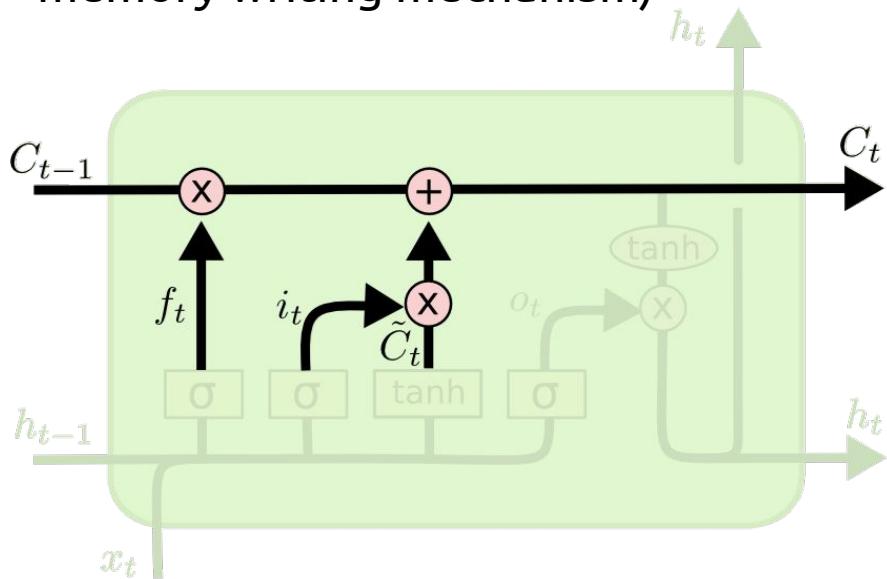
New contribution to cell state

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Classic neuron

Long Short-Term Memory (LSTM)

Make every RNN unit able to decide how to **to update the cell state** (optimized memory writing mechanism)

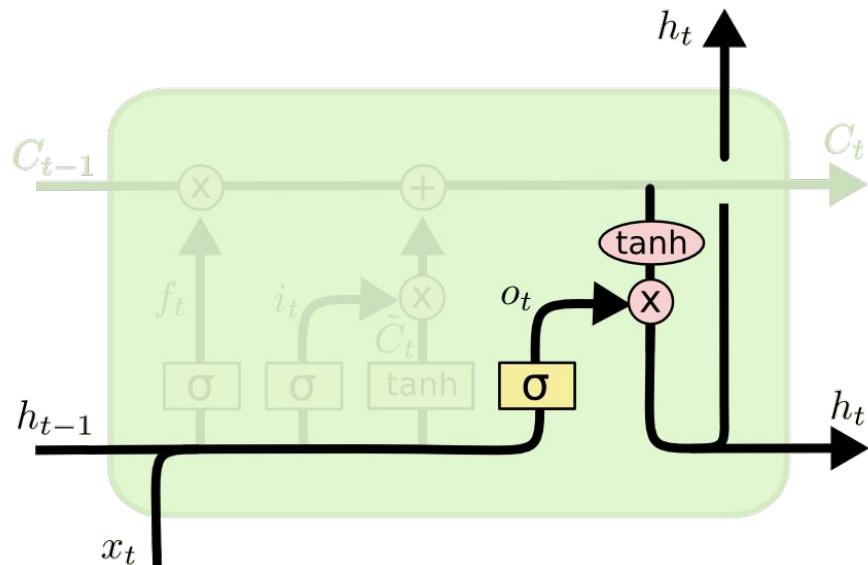


**Forget + Input Gates =
Update Cell State (memory):**

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Long Short-Term Memory (LSTM)

Make every RNN unit able to **output the decisions whenever it is ready to do so** (optimized output mechanism)



Output Gate Layer

$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

Output to next layer & timestep

$$h_t = o_t * \tanh (C_t)$$

Long Short-Term Memory (LSTM)

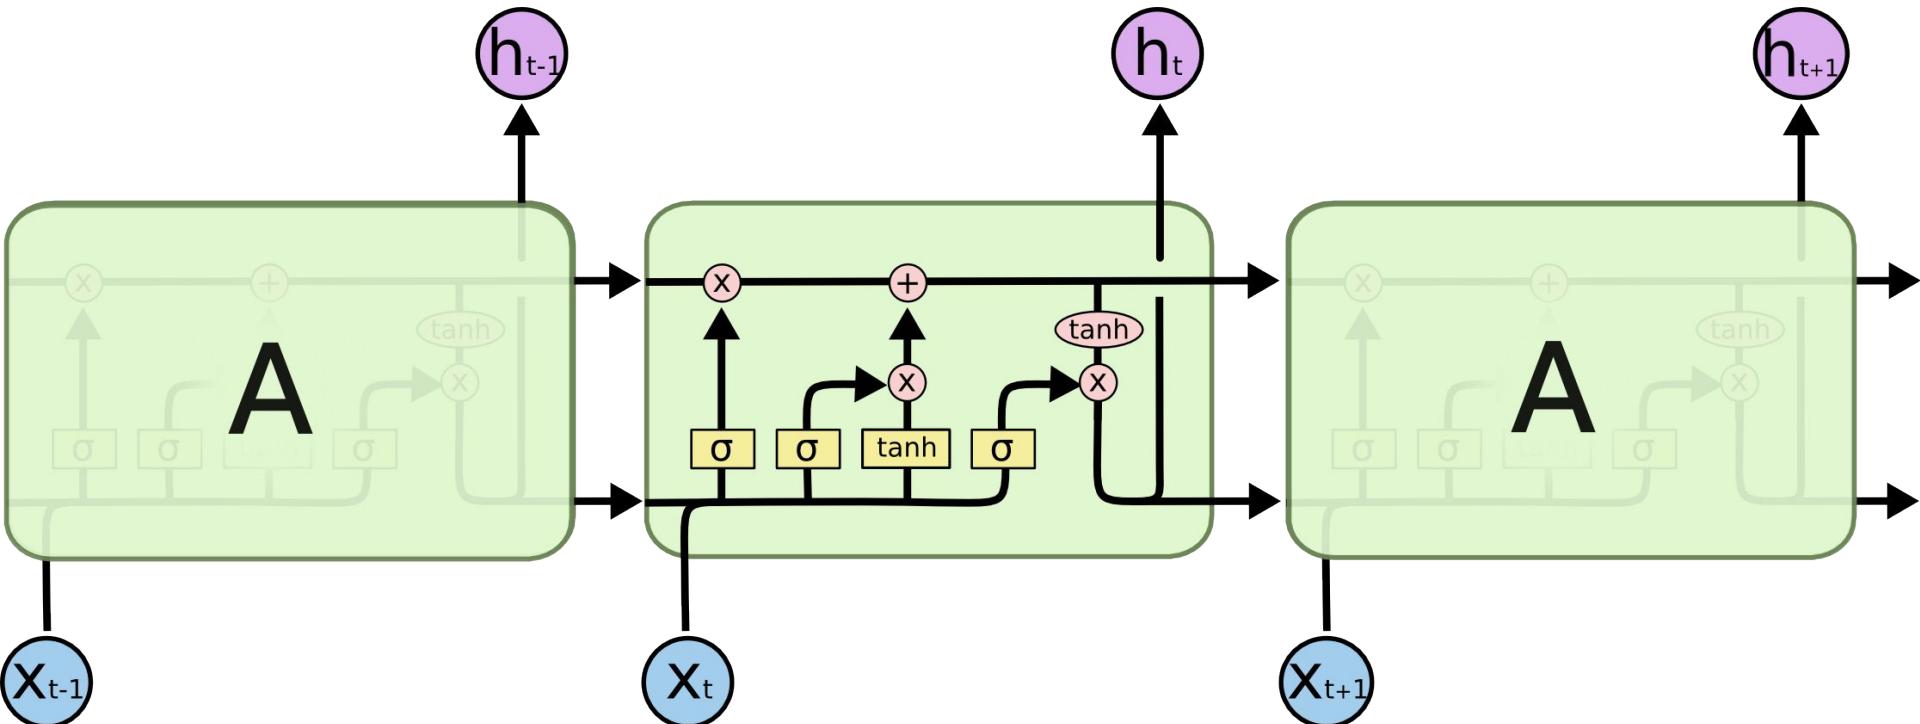
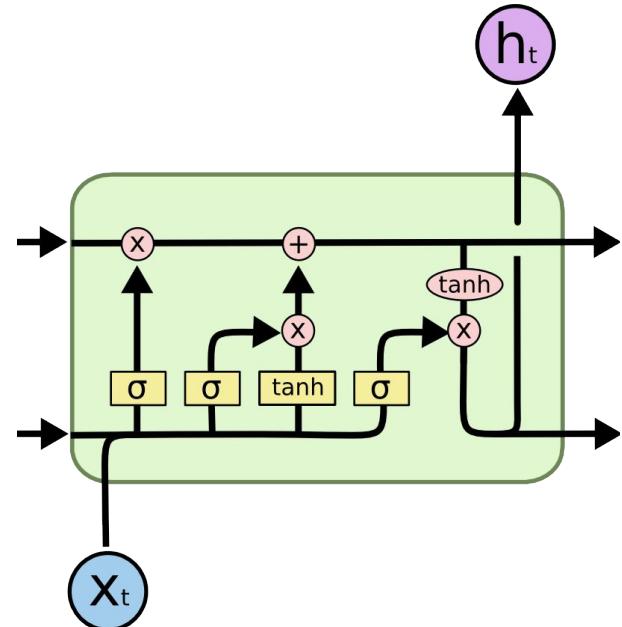


Figure: Cristopher Olah, ["Understanding LSTM Networks"](#) (2015)

Long Short-Term Memory (LSTM)



Compared to a non-gated RNN, an LSTM has four times more parameters because of the additional neurons that govern the gates:

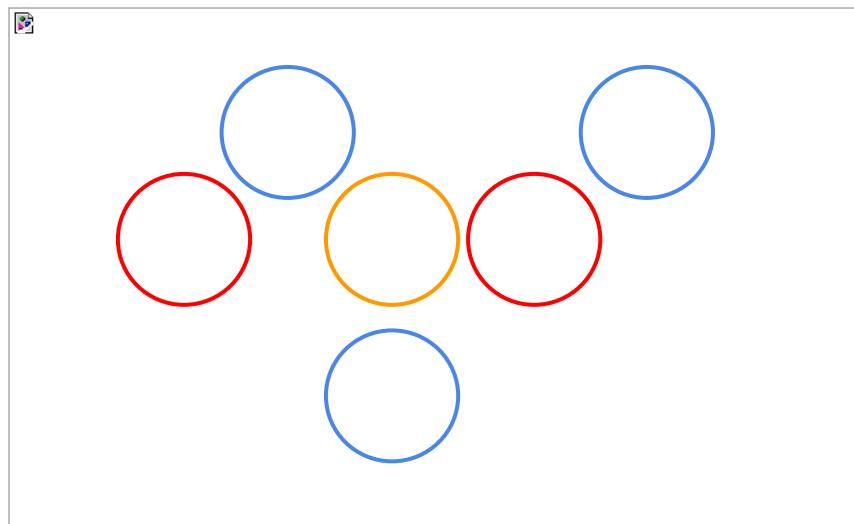
$$N_{params}^i = 4 \times (N_{inputs}^i \times N_{units}^i + N_{units}^i \times N_{units}^i + N_{units}^i)$$

3 sigmoid gates +
input activation (\tanh in the figure)

Long Short Term Memory (LSTM) cell

Updating an LSTM cell requires 6 computations:

1. Gates
2. Activation units
3. Cell state



Computation Flow

$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$

$\hat{C}_t = \tanh(W_c x_t + U_c h_{t-1} + b_c)$

$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$

$C_t = i_t \odot \hat{C}_t + f_t \odot C_{t-1}$

$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$

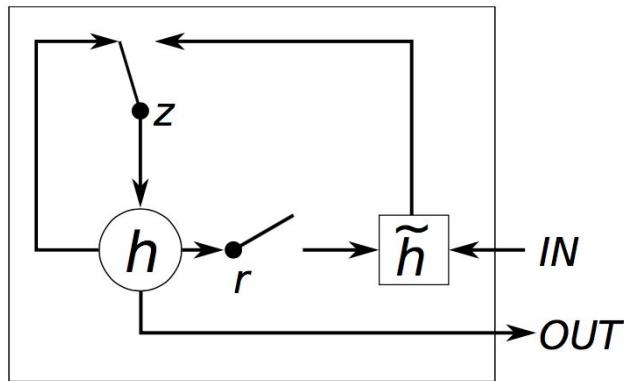
$h_t = o_t \odot \tanh(C_t)$

Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time
4. Gated Recurrent Neurons
 - a. LSTM
 - b. GRU

Gated Recurrent Unit (GRU)

GRU obtain a similar performance as LSTM with one gate less.



$$u_i = \sigma(W^{(u)}x_i + U^{(u)}h_{i-1} + b^{(u)}) \quad (1)$$

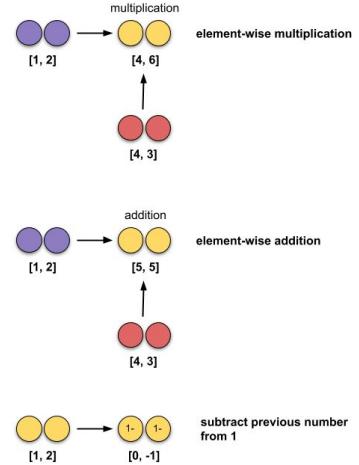
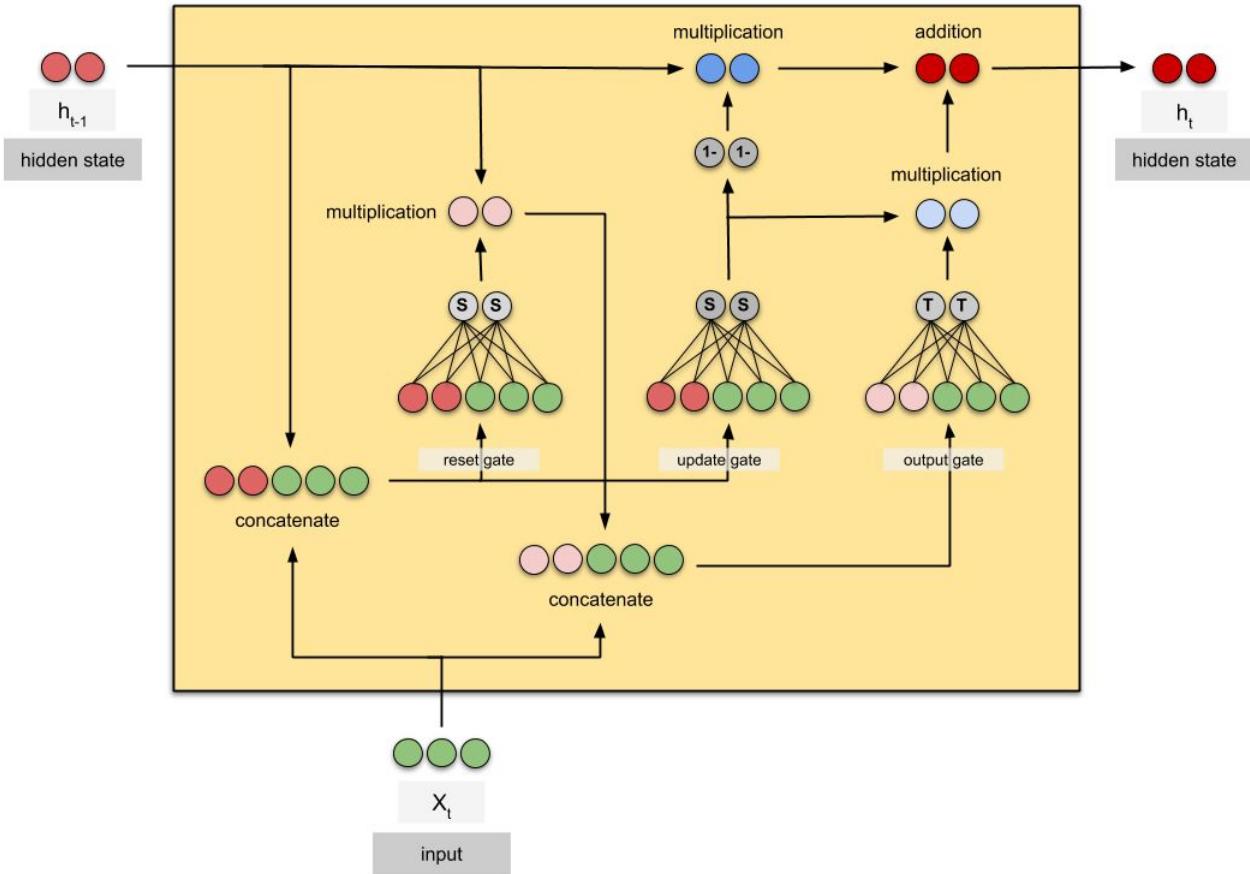
$$r_i = \sigma(W^{(r)}x_i + U^{(r)}h_{i-1} + b^{(r)}) \quad (2)$$

$$\tilde{h}_i = \tanh(Wx_i + r_i \circ Uh_{i-1} + b^{(h)}) \quad (3)$$

$$h_i = u_i \circ \tilde{h}_i + (1 - u_i) \circ h_{i-1} \quad (4)$$

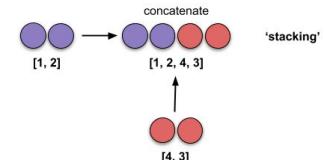
$$N_{params}^i = 3 \times (N_{inputs}^i \times N_{units}^i + N_{units}^i \times N_{units}^i + N_{units}^i)$$

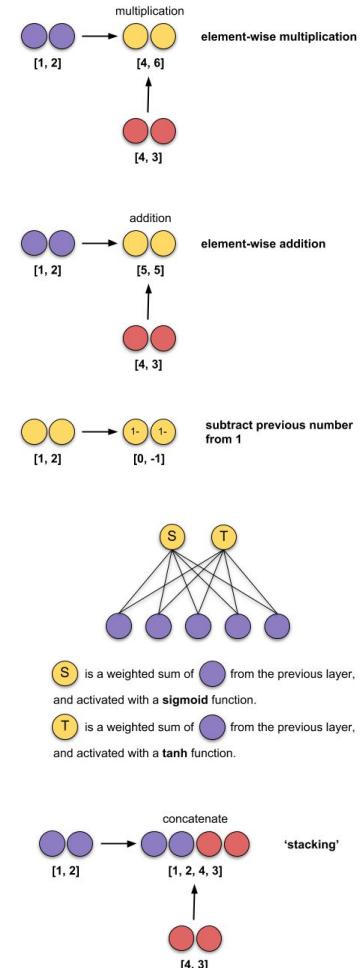
#GRU Cho, Kyunghyun, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. ["Learning phrase representations using RNN encoder-decoder for statistical machine translation."](#) EMNLP 2014.



S is a weighted sum of h_{t-1} from the previous layer and activated with a sigmoid function.

T is a weighted sum of x_t from the previous layer and activated with a tanh function.





Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time
4. Gated Recurrent Neurons
 - a. LSTM
 - b. GRU
5. Stacked RNNs

Deep (stacked) RNNs

Recurrent layers can be stacked as any non-recurrent layer.

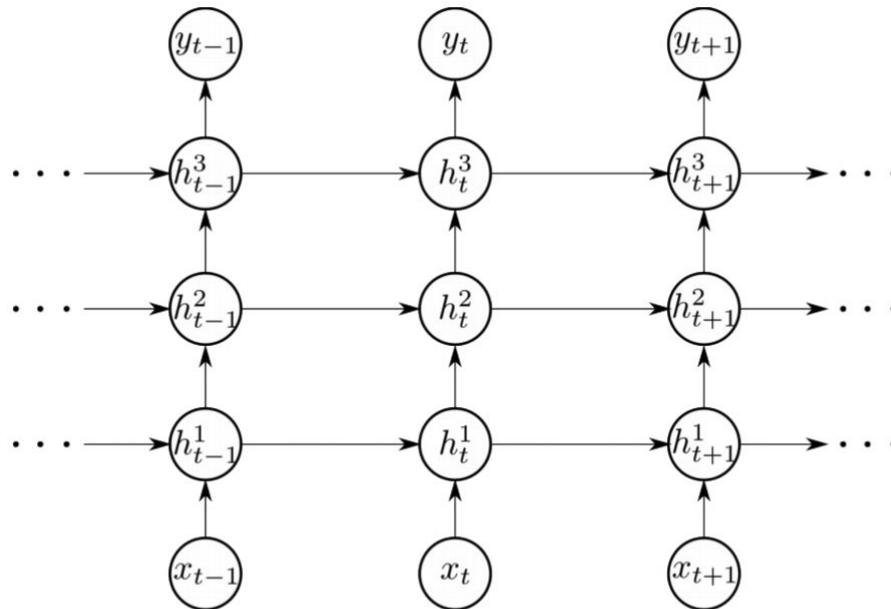


Figure:
[Lambert](#)
(Stanford University).

Bidirectional RNN (BRNN)

Sequences can be process in both directions by stacking two RNN layers.

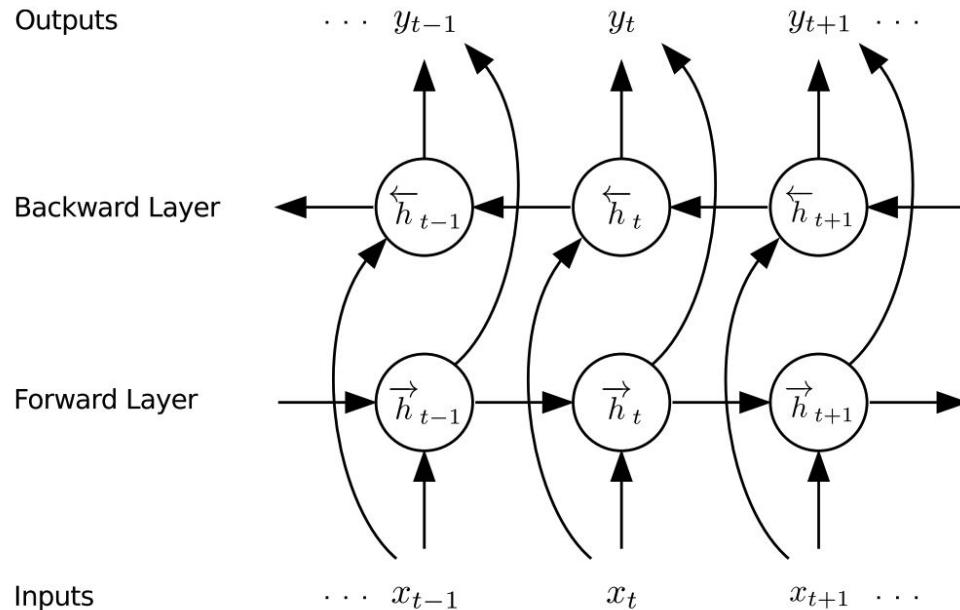


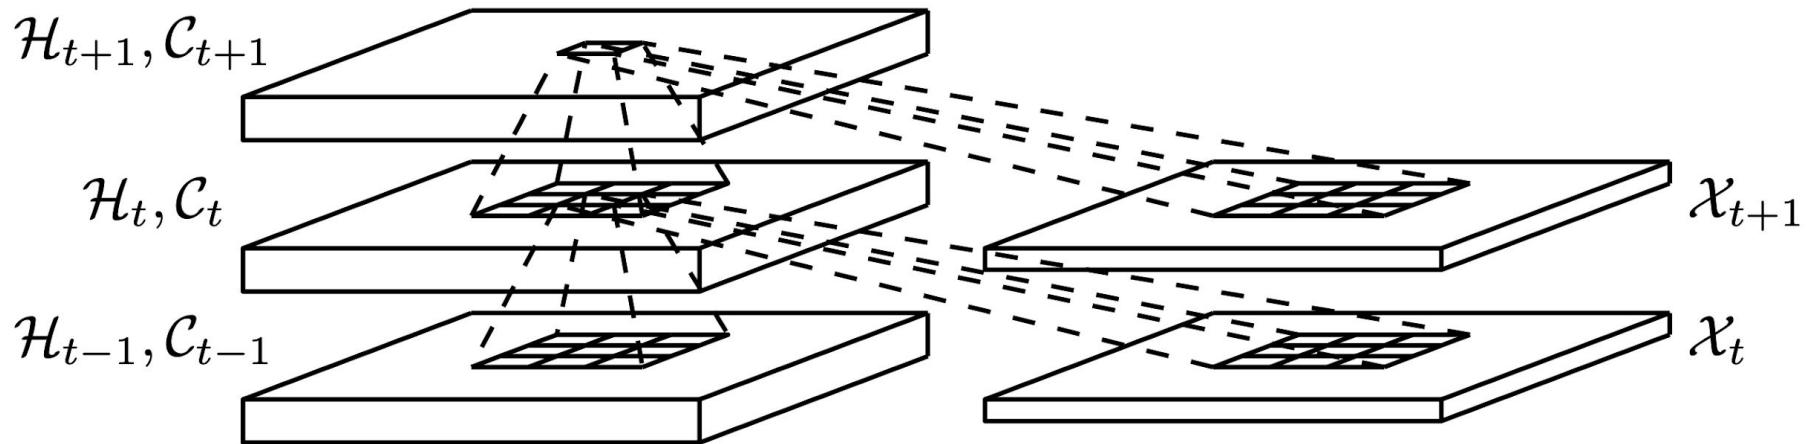
Figure:
Graves et al. 2013.

#BRNN Schuster, M., & Paliwal, K. K. (1997). [Bidirectional recurrent neural networks](#). IEEE transactions on Signal Processing, 45(11), 2673-2681.

Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time
4. Gated Recurrent Neurons
 - a. LSTM
 - b. GRU
5. Stacked RNNs
6. Advanced architectures

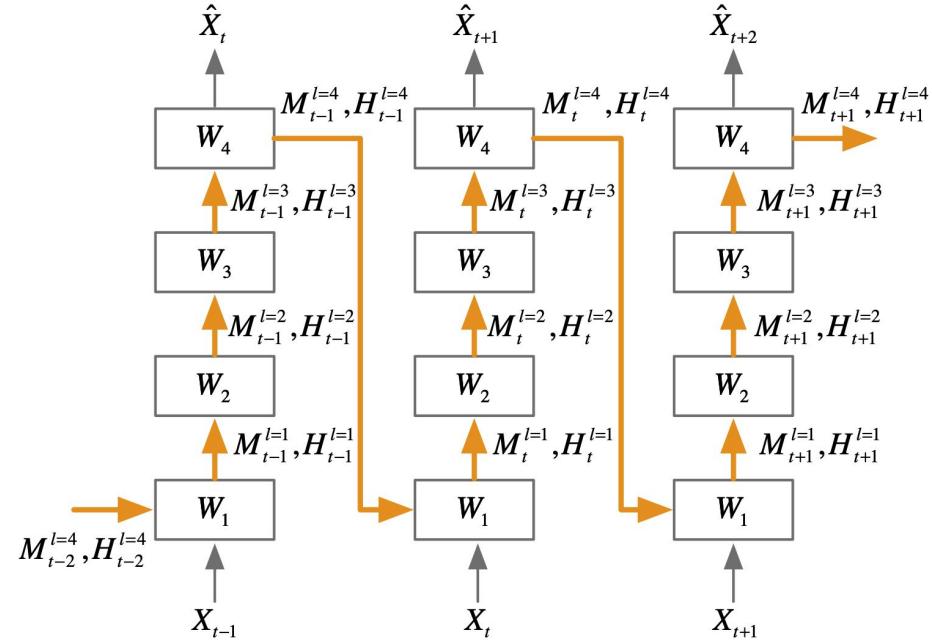
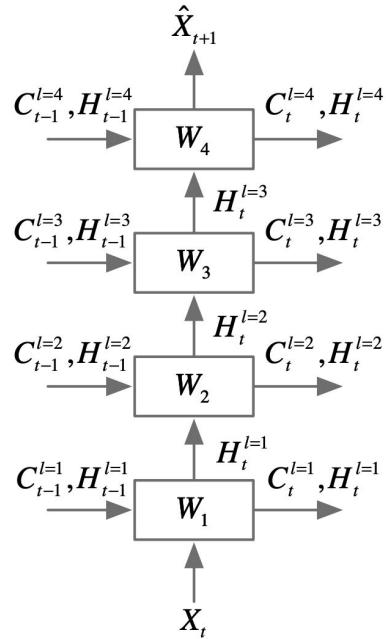
Conv-LSTM



#ConvLSTM Xingjian, S. H. I., Chen, Z., Wang, H., Yeung, D. Y., Wong, W. K., & Woo, W. C. [Convolutional LSTM network: A machine learning approach for precipitation nowcasting](#). NIPS 2015.

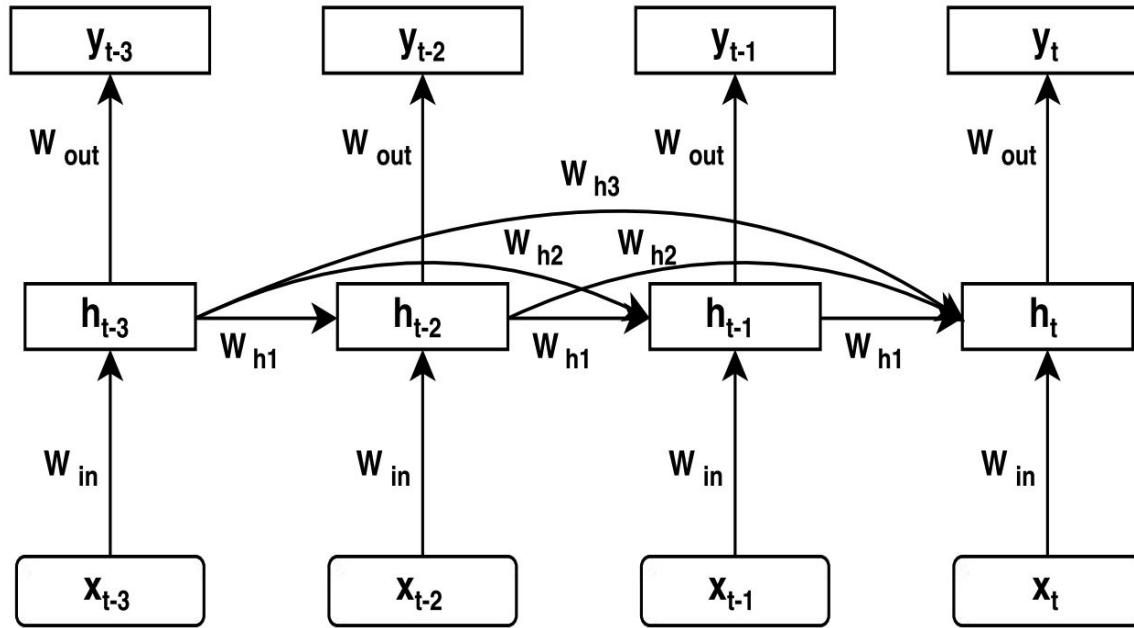
Spatiotemporal LSTM (ST-LSTM)

Cell states are no longer constrained inside each LSTM unit. Instead, they are allowed to zigzag in two directions: across stacked RNN layers vertically and through all RNN states horizontally

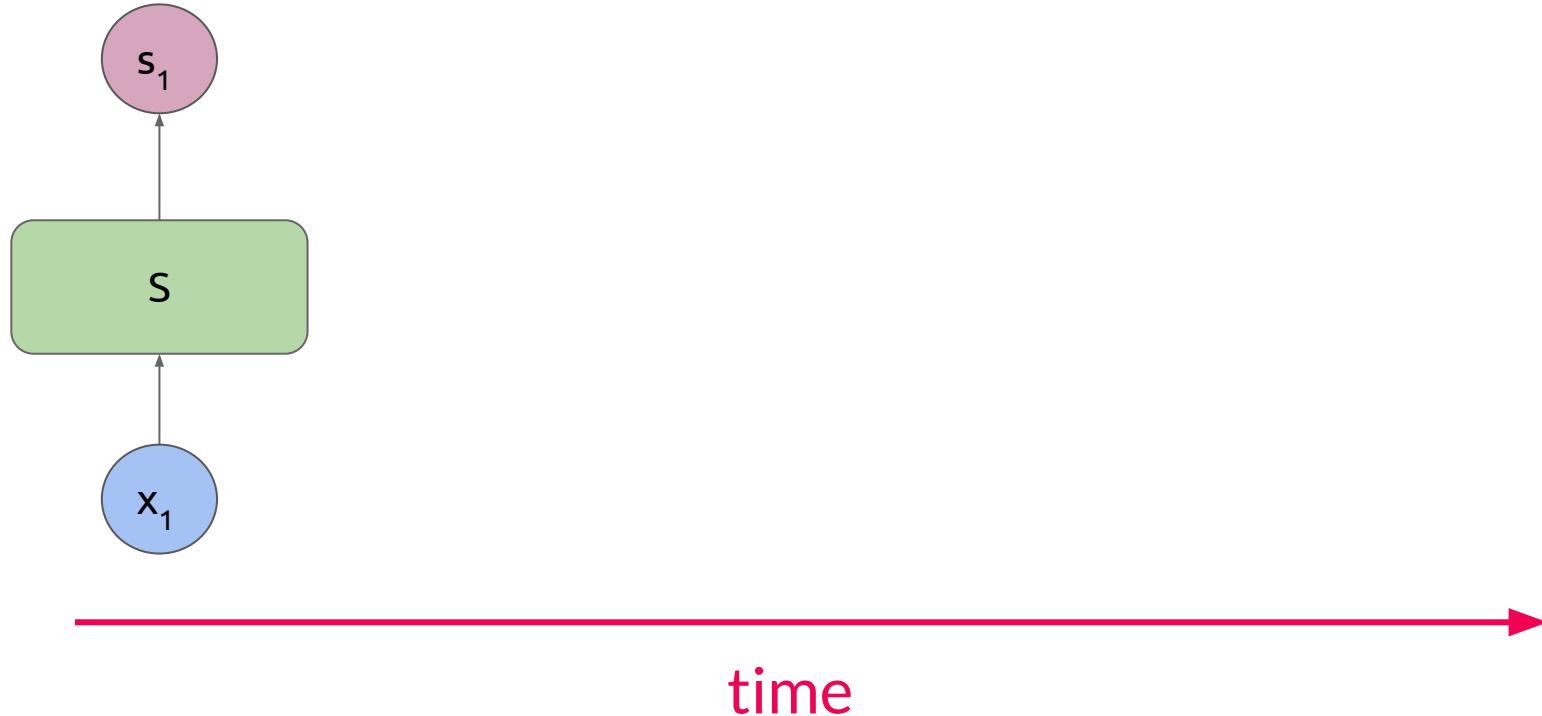


Higher Order RNN

More memory units can be added to connect with further temporal states.

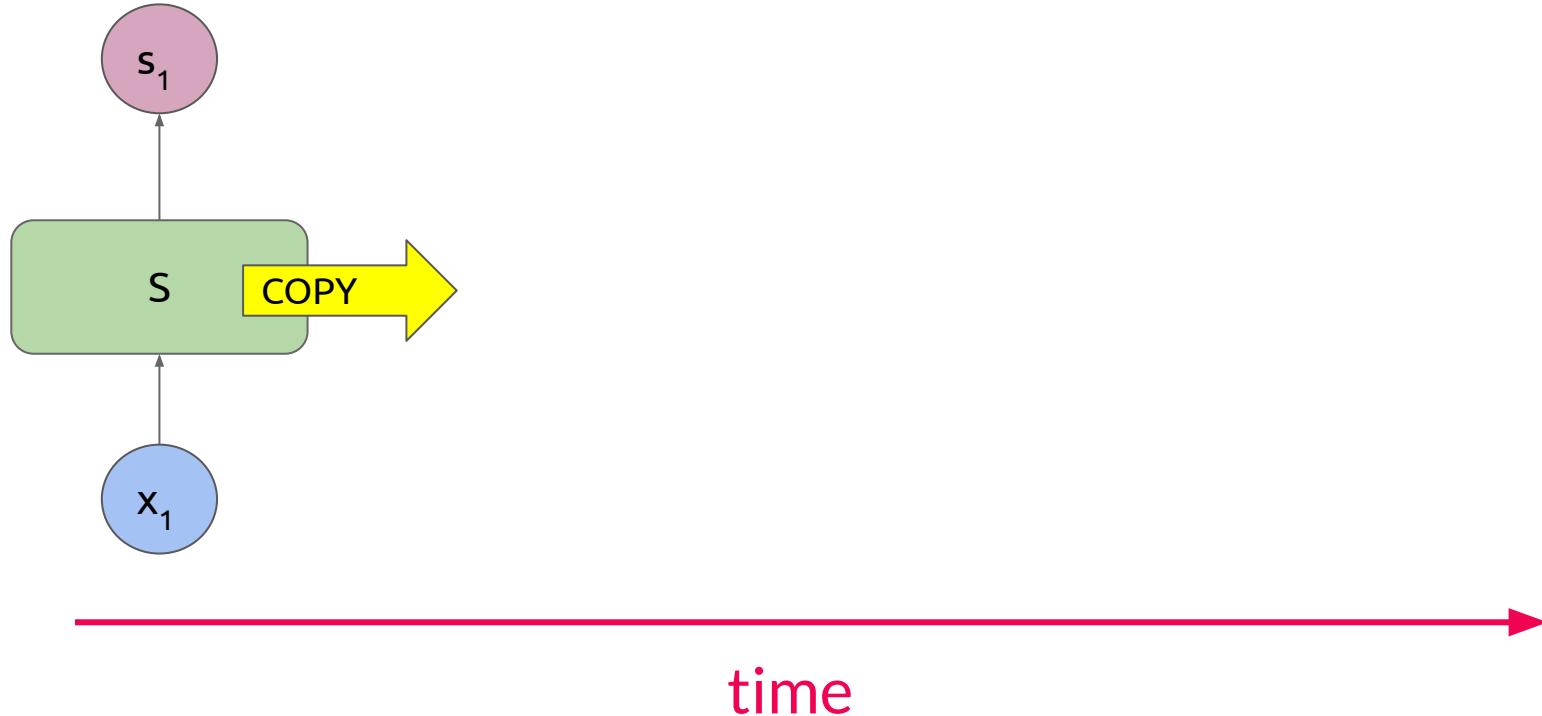


SkipRNN



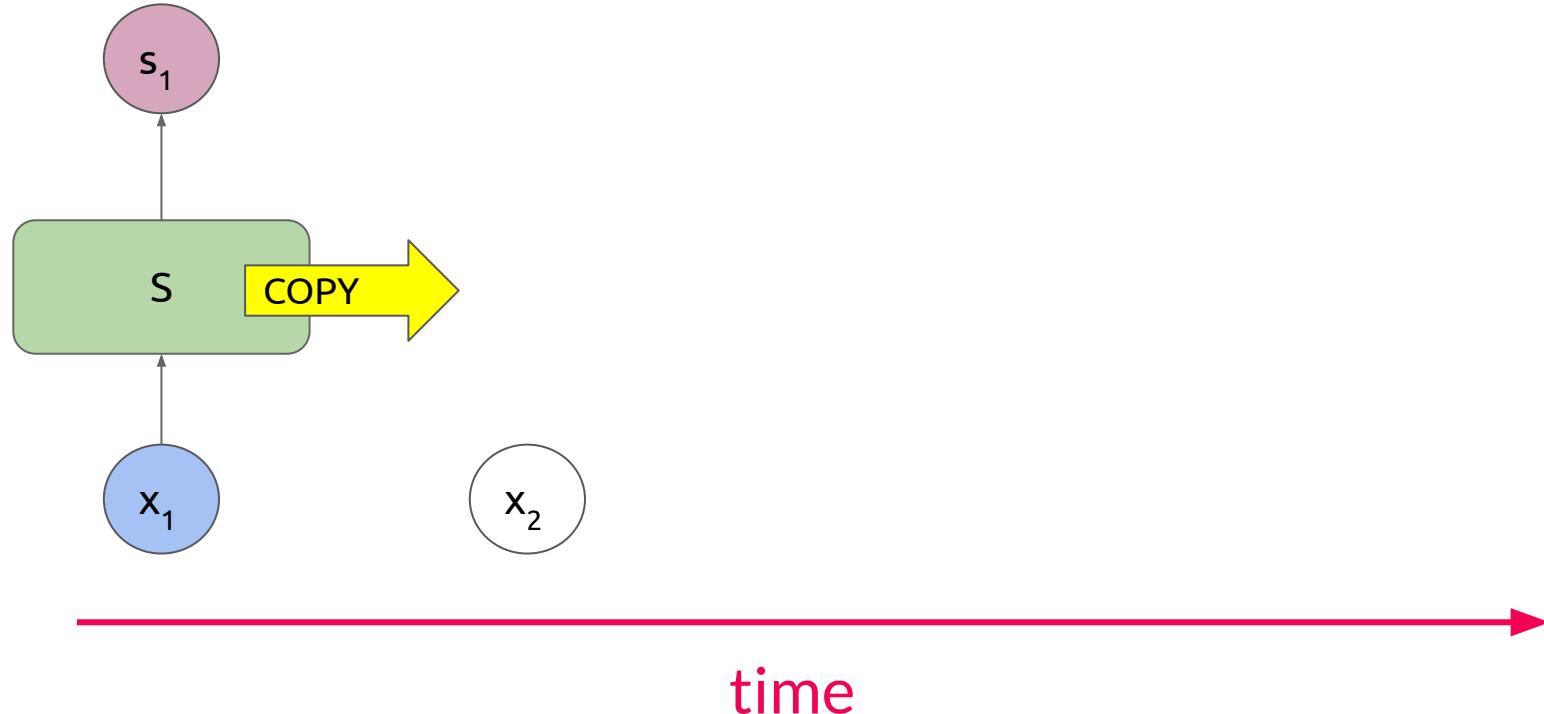
#SkipRNN Victor Campos, Brendan Jou, Xavier Giro-i-Nieto, Jordi Torres, and Shih-Fu Chang. ["Skip RNN: Learning to Skip State Updates in Recurrent Neural Networks"](#), ICLR 2018.

SkipRNN



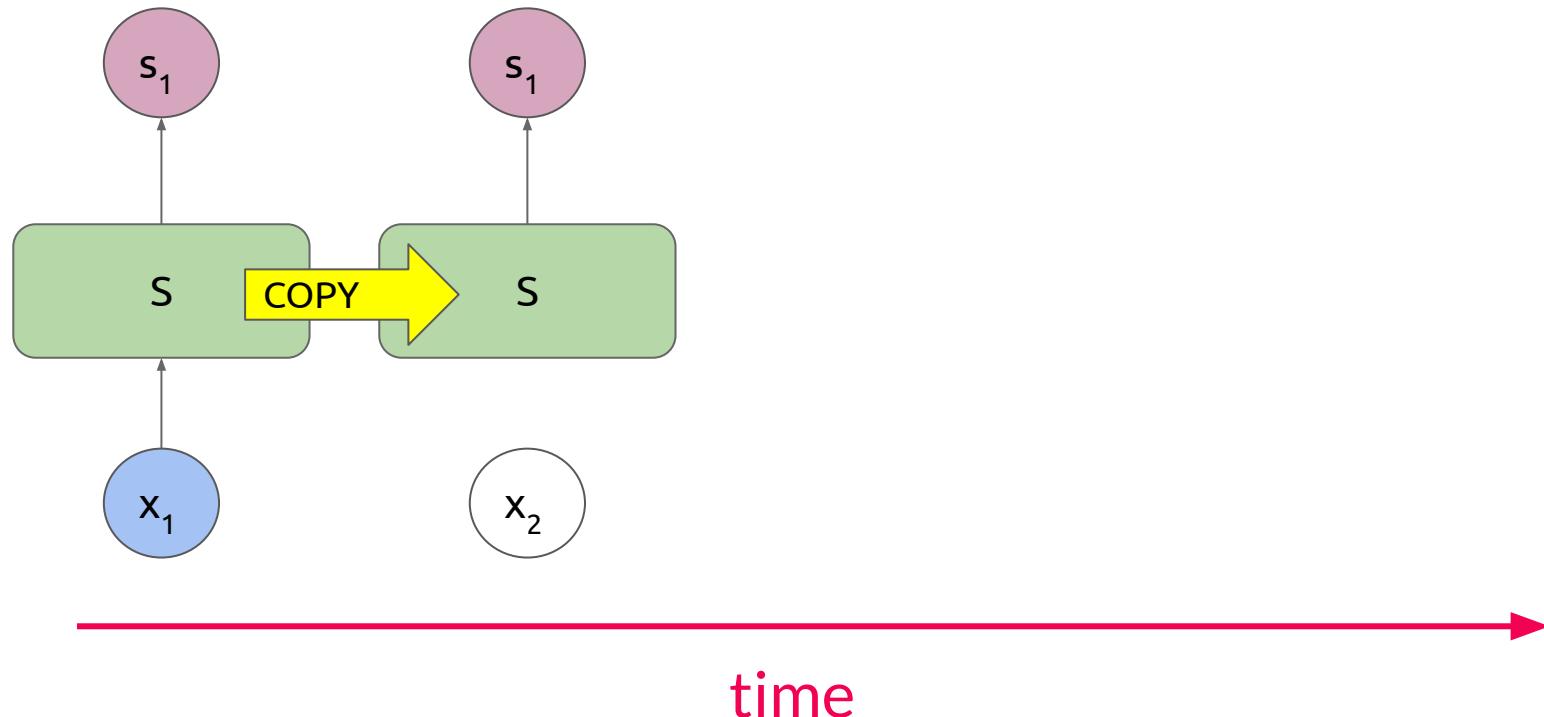
#SkipRNN Victor Campos, Brendan Jou, Xavier Giro-i-Nieto, Jordi Torres, and Shih-Fu Chang. ["Skip RNN: Learning to Skip State Updates in Recurrent Neural Networks"](#), ICLR 2018.

SkipRNN



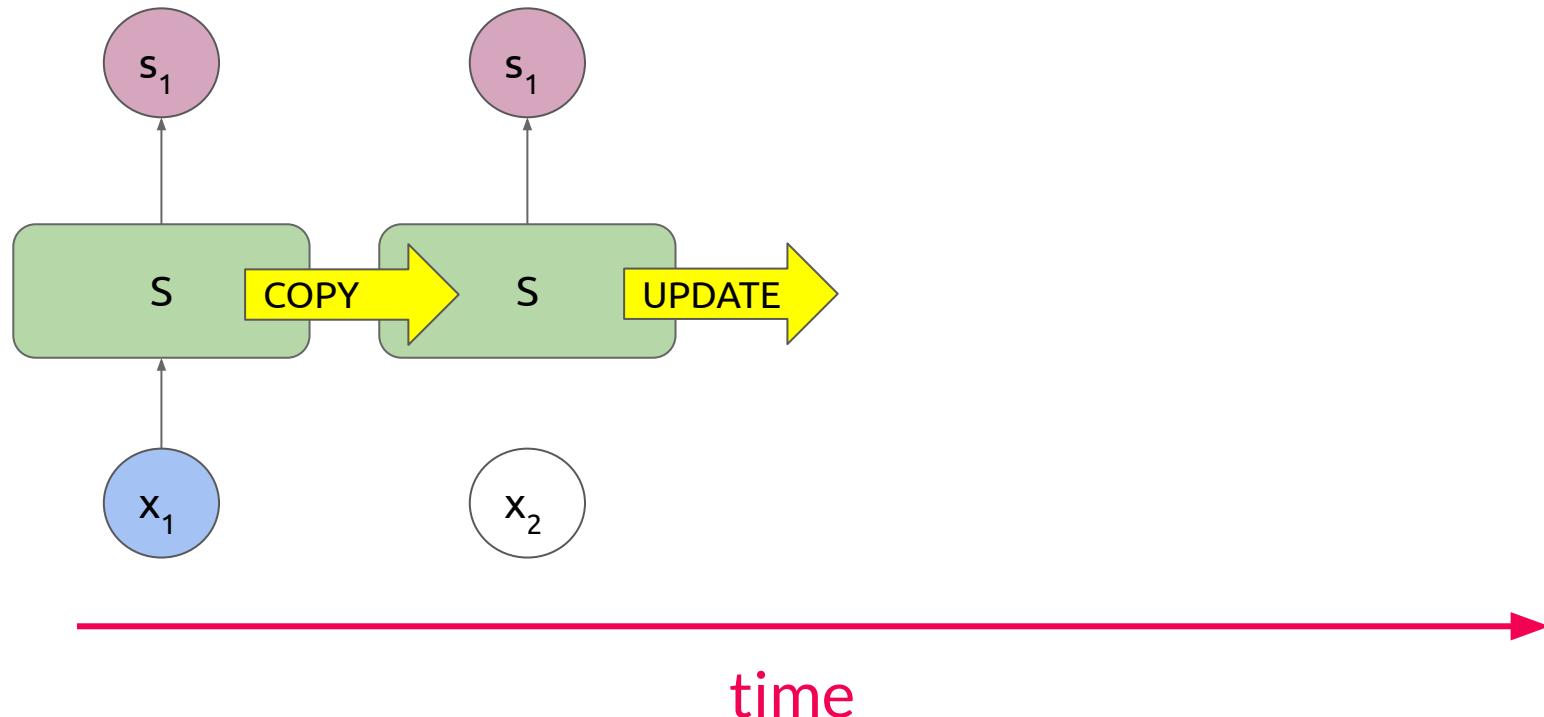
#SkipRNN Victor Campos, Brendan Jou, Xavier Giro-i-Nieto, Jordi Torres, and Shih-Fu Chang. ["Skip RNN: Learning to Skip State Updates in Recurrent Neural Networks"](#), ICLR 2018.

SkipRNN



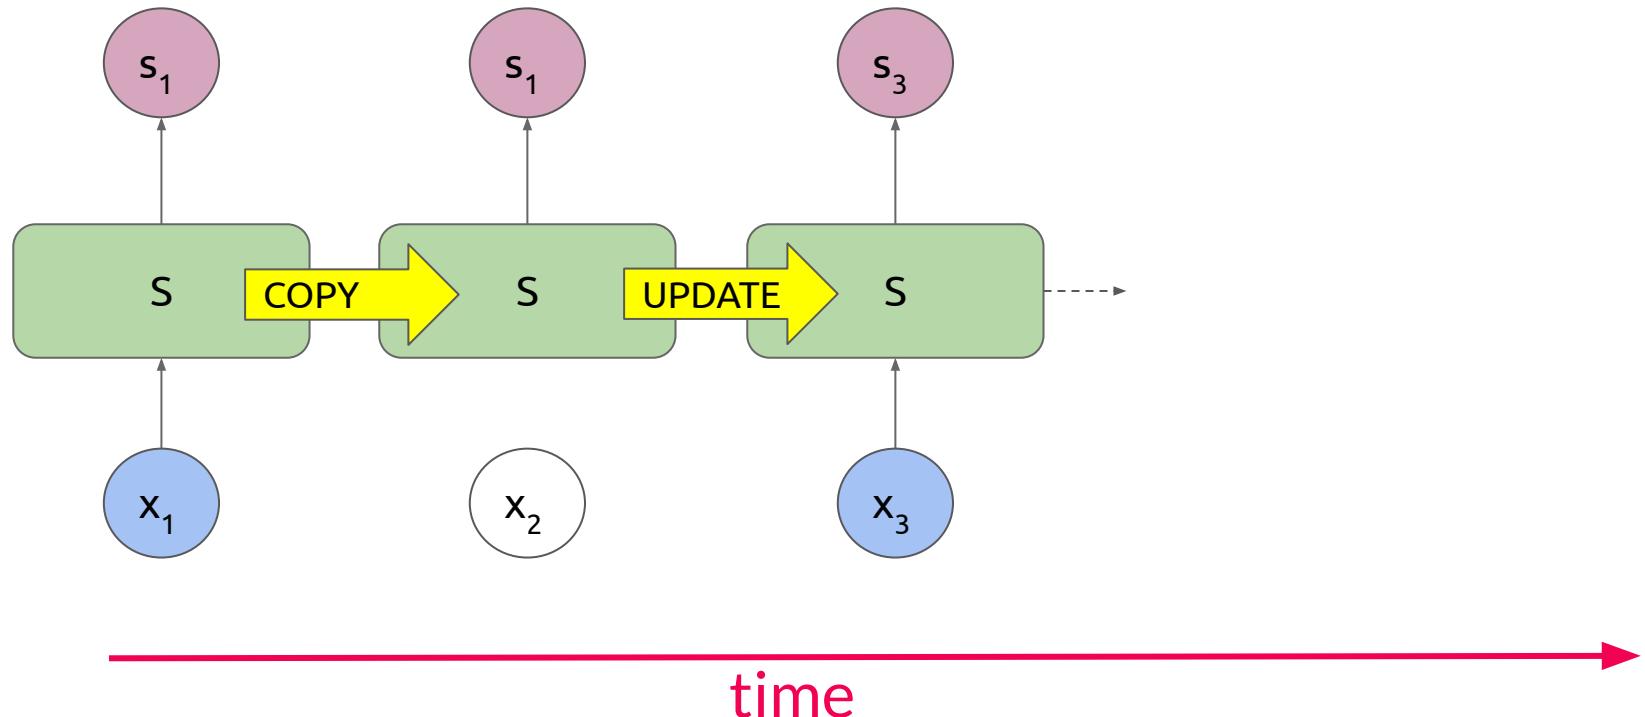
#SkipRNN Victor Campos, Brendan Jou, Xavier Giro-i-Nieto, Jordi Torres, and Shih-Fu Chang. ["Skip RNN: Learning to Skip State Updates in Recurrent Neural Networks"](#), ICLR 2018.

SkipRNN



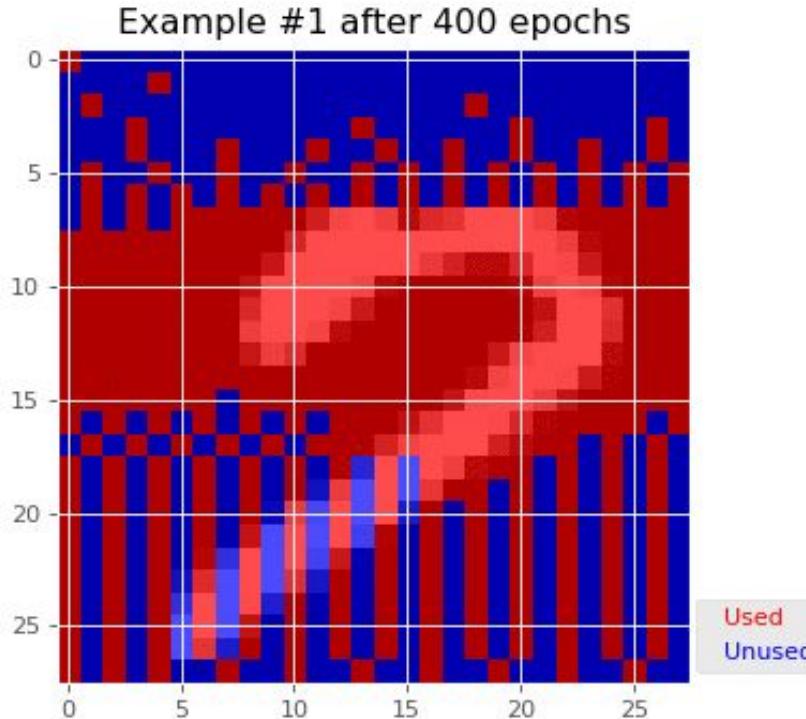
#SkipRNN Victor Campos, Brendan Jou, Xavier Giro-i-Nieto, Jordi Torres, and Shih-Fu Chang. ["Skip RNN: Learning to Skip State Updates in Recurrent Neural Networks"](#), ICLR 2018.

SkipRNN



#SkipRNN Victor Campos, Brendan Jou, Xavier Giro-i-Nieto, Jordi Torres, and Shih-Fu Chang. ["Skip RNN: Learning to Skip State Updates in Recurrent Neural Networks"](#), ICLR 2018.

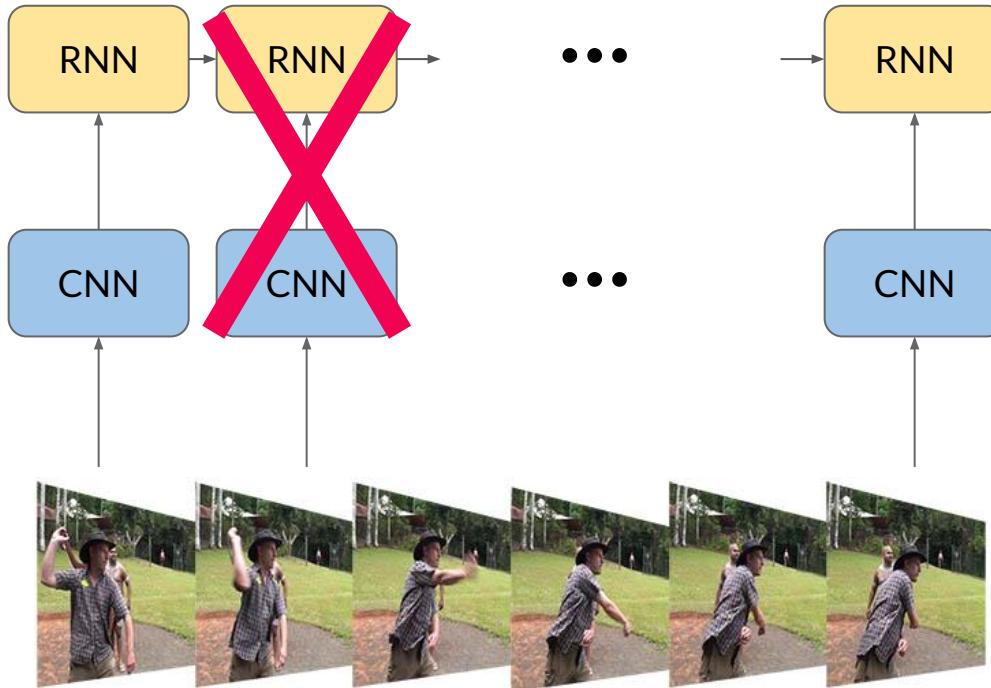
SkipRNN



~95% acc

Used
Unused

SkipRNN



After processing a frame, let the RNN decide how many future frames can be skipped

In skipped frames, simply copy the output and state from the previous time step

There is no ground truth for which frames can be skipped. The RNN **learns** it by itself during training!

SkipRNN

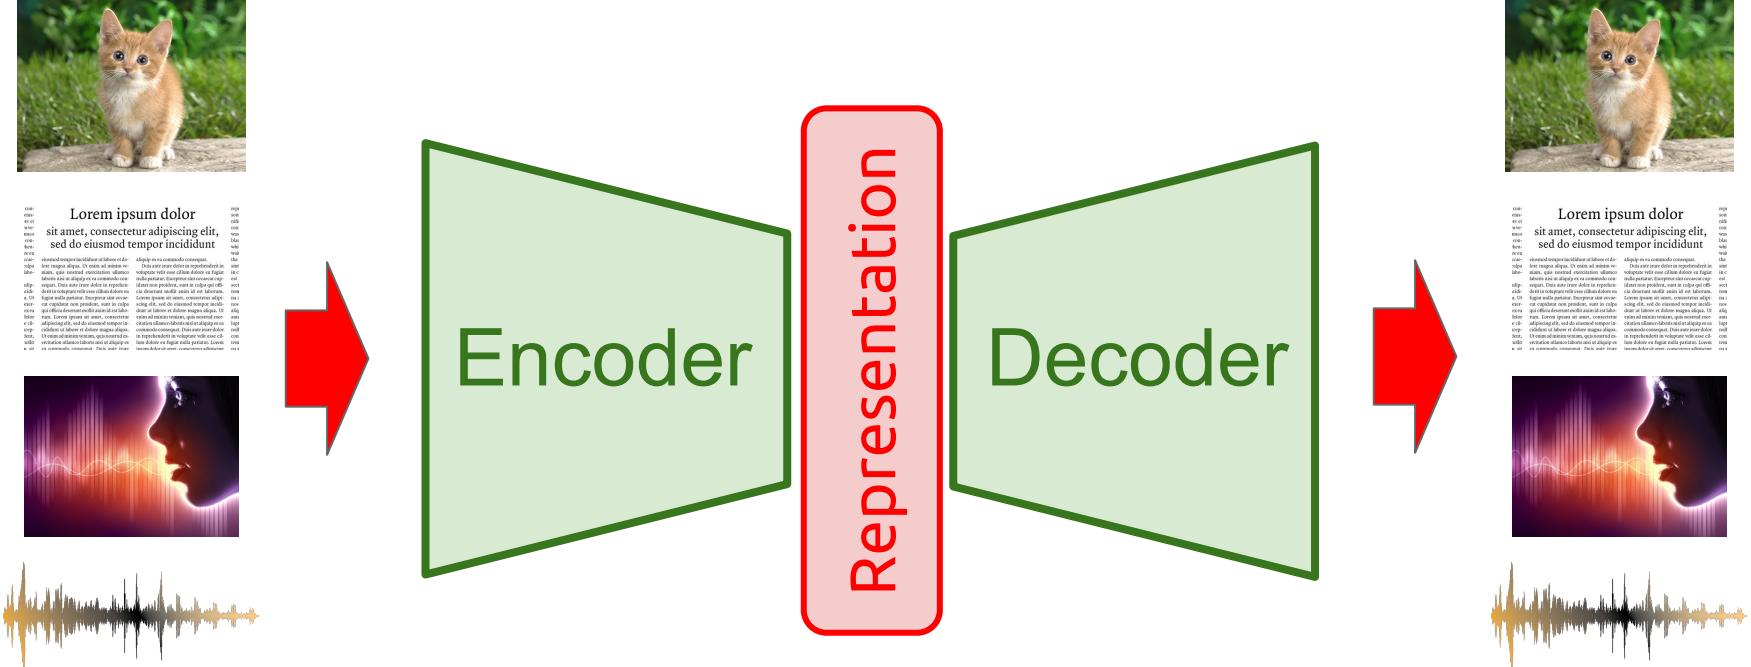


Used
Unused

Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time
4. Gated Recurrent Neurons
 - a. LSTM
 - b. GRU
5. Advanced Architectures
6. Encoder-Decoder Architectures

Encoder-Decoder Applications



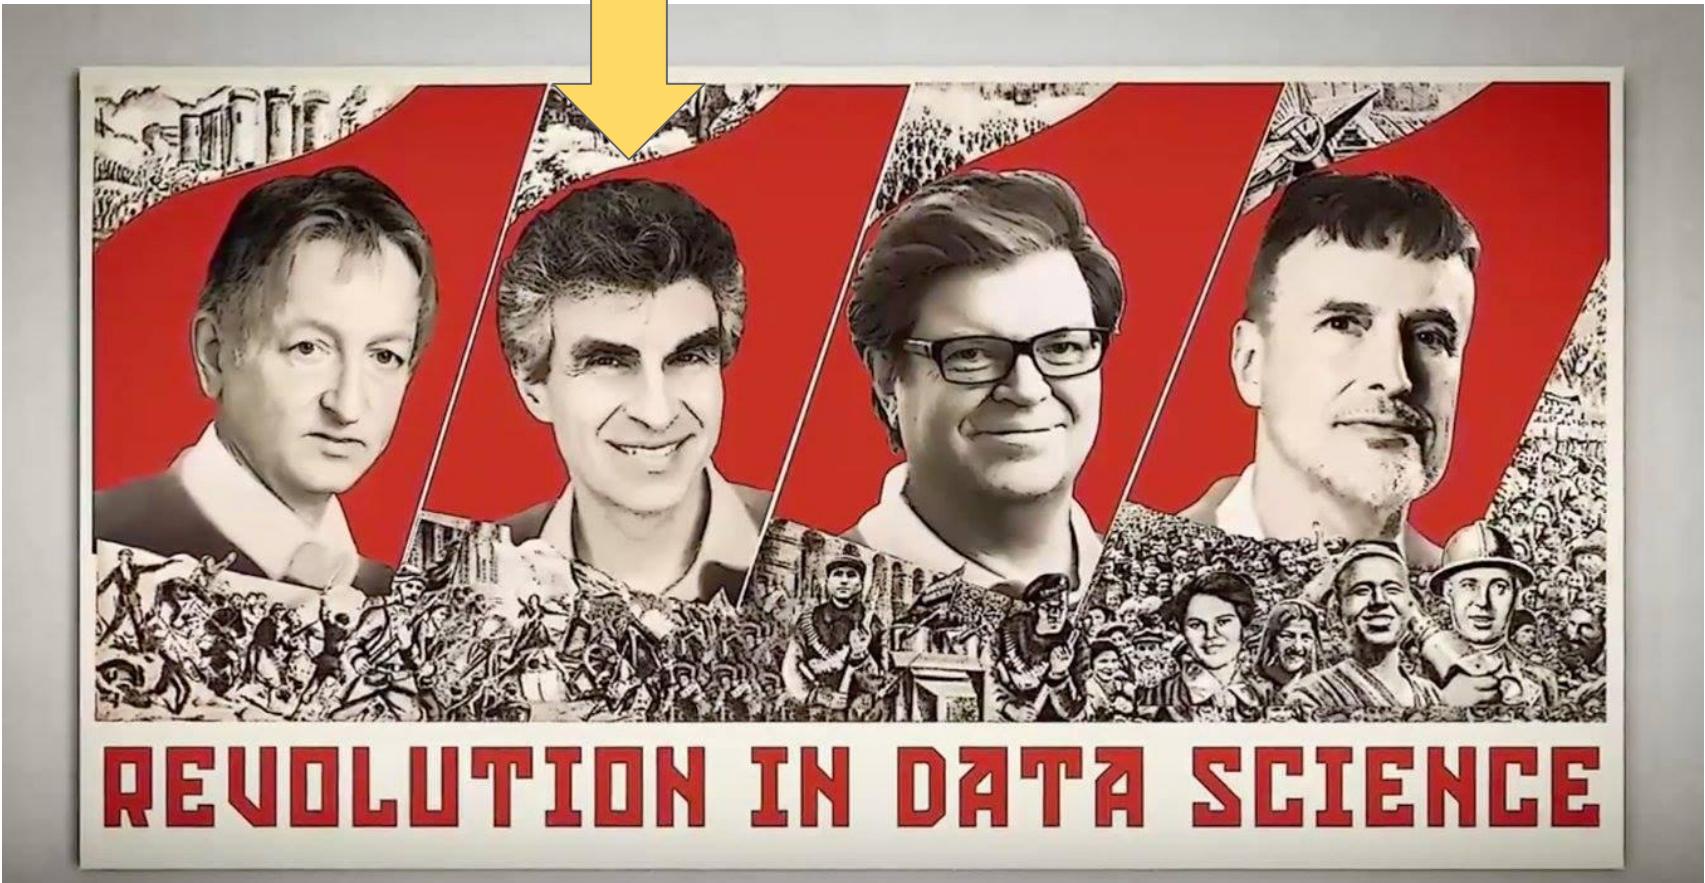


Figure: Michael Bronstein, "[Geometric Deep Learning](#)" (ICLR 2021)

Encoding of a Sequence of Words

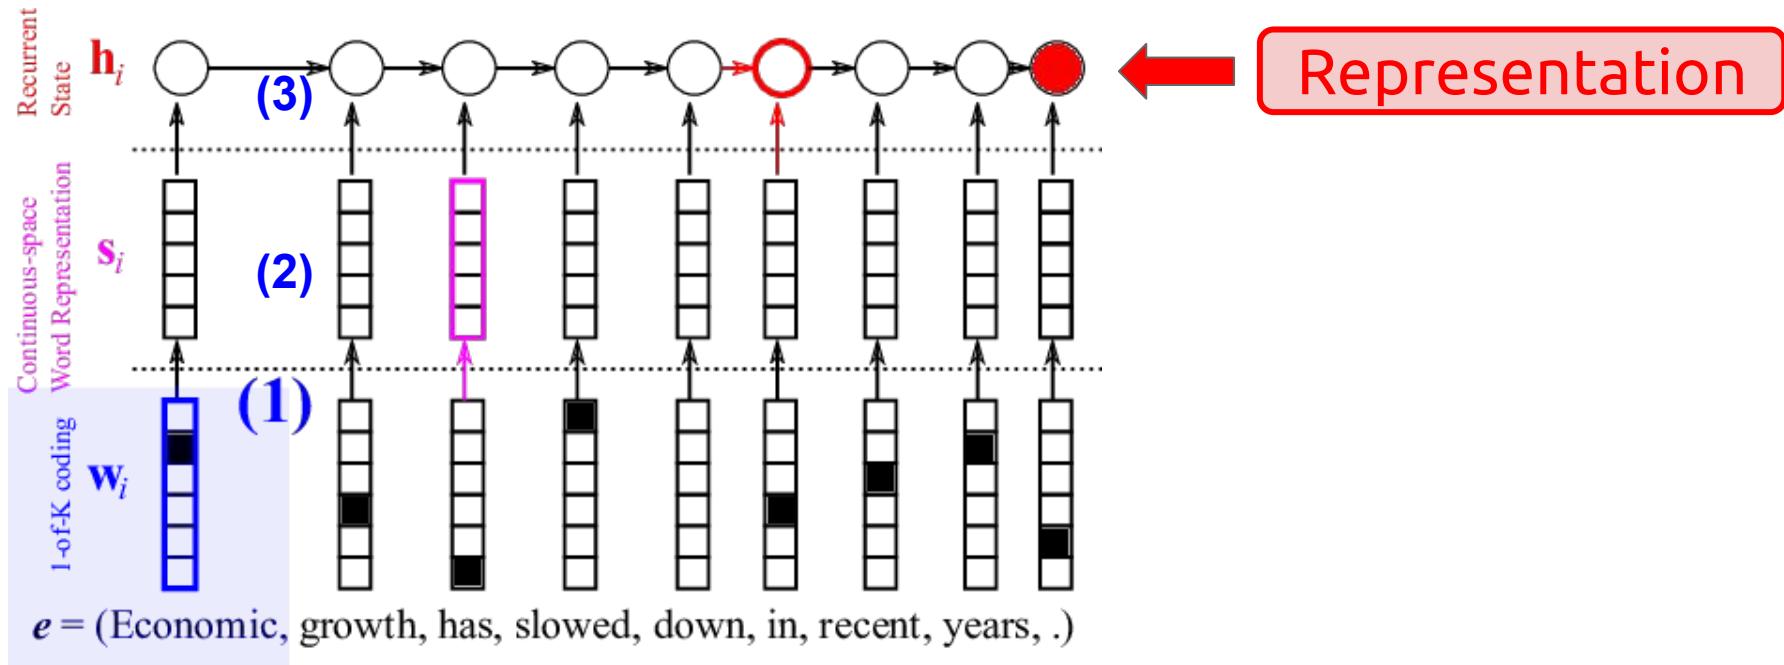
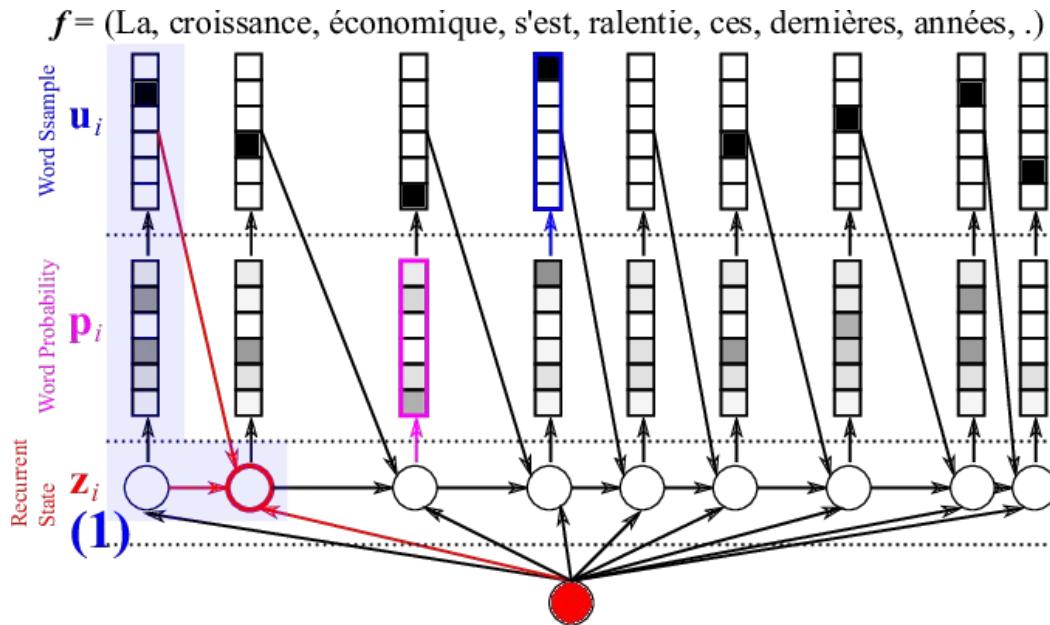


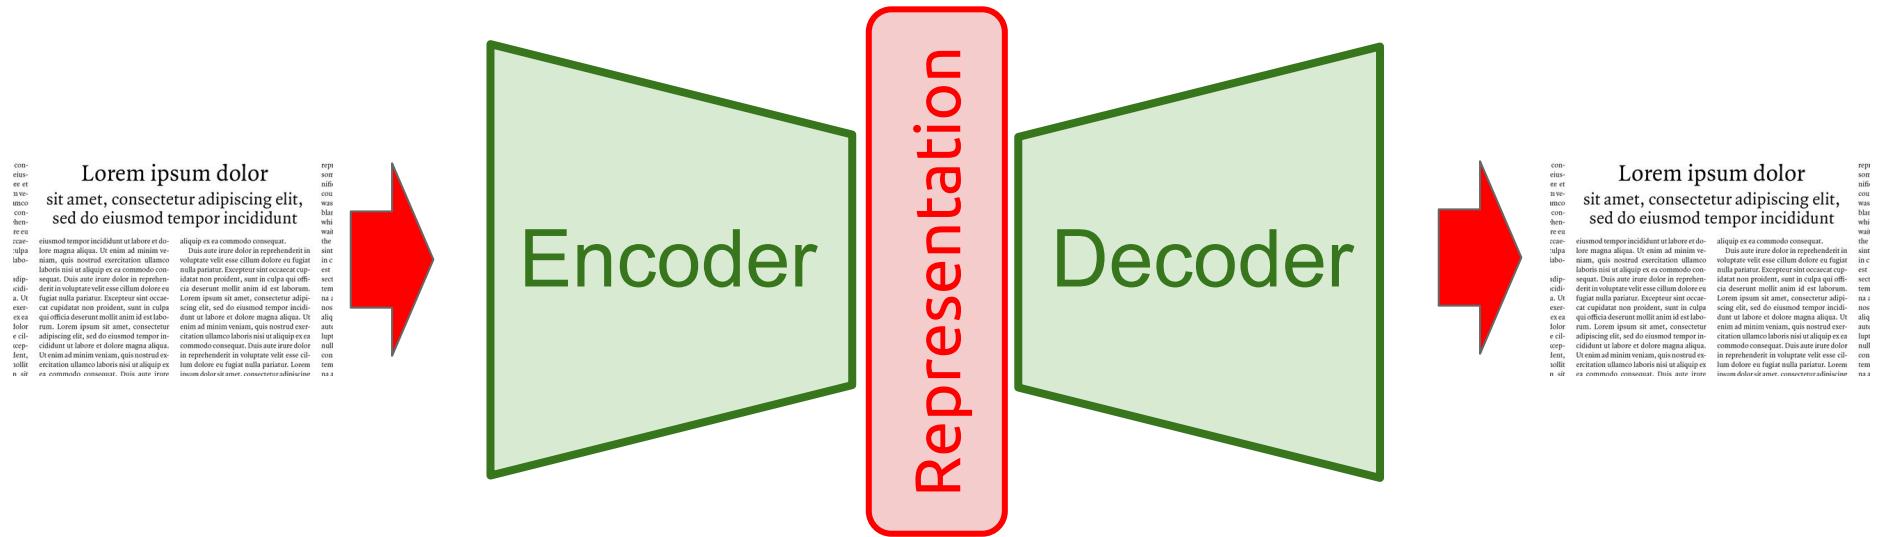
Fig: Kyunghyun Cho, "[Introduction to Neural Machine Translation with GPUs](#)" (2015)

Cho, Kyunghyun, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and **Yoshua Bengio**. "[Learning phrase representations using RNN encoder-decoder for statistical machine translation.](#)" EMNLP 2014.

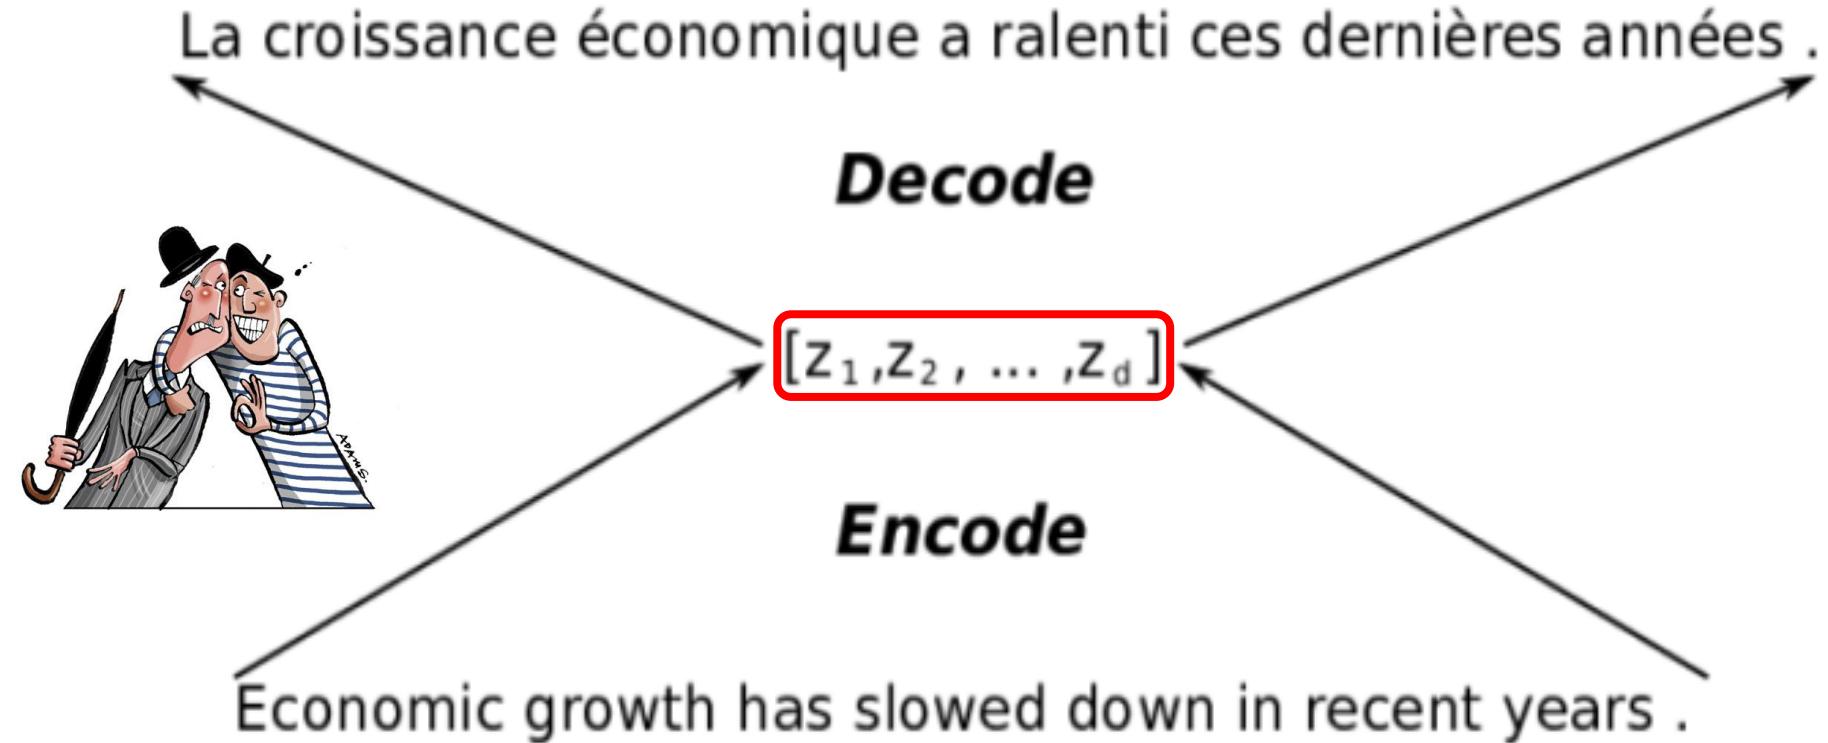
Decoding of a Sequence of Words



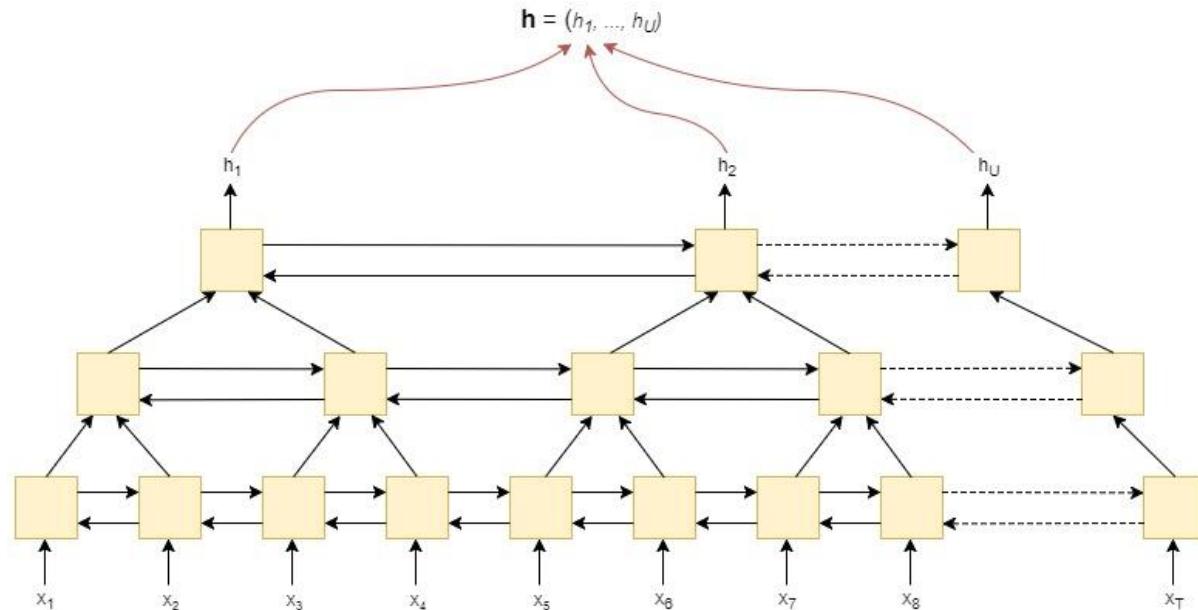
Representation



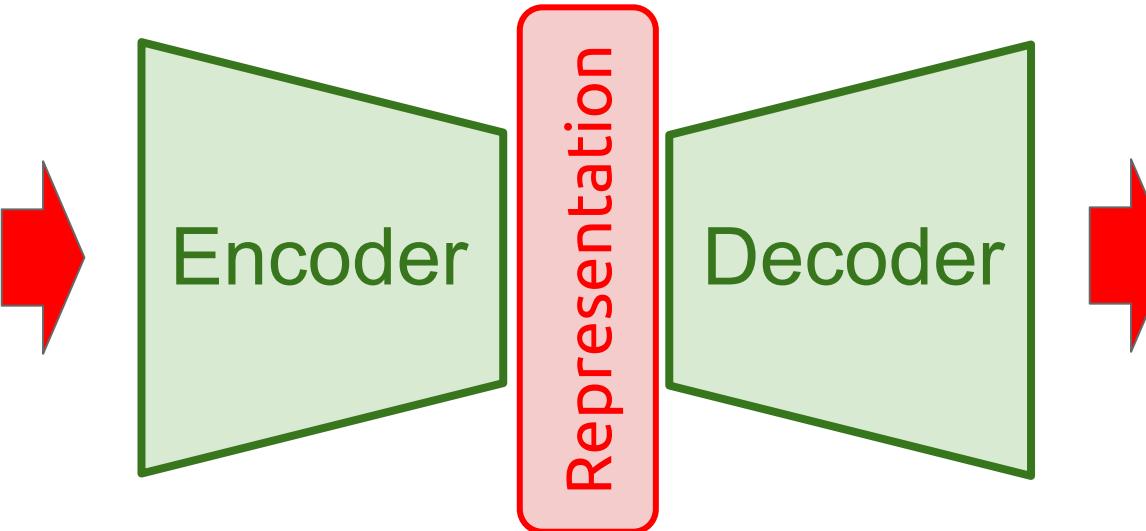
Neural Machine Translation (NMT)



Speech Encoding



Chan, William, Navdeep Jaitly, Quoc Le, and Oriol Vinyals. ["Listen, attend and spell: A neural network for large vocabulary conversational speech recognition."](#) ICASSP 2016.



elusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

Duis autem irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Jt fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa
e- Lorem ipsum sit amet, consectetur adipisci-
ng elit, sed do eiusmod tempore incidi-

qui officia deserunt mollit anim id est laborum. Lorem ipsum sit amet, consectetur

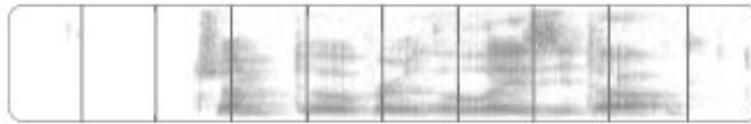
etiam ad minimam remanit, quis noscitur ex actione ullamco laboris nisi ut aliquip ex ea.

cididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud ex- commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cil-

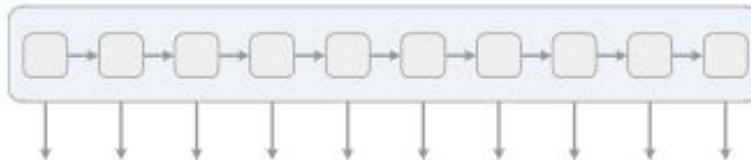
ercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

ea commodo concedit. Tunc ante irum inum dolor sat amet, conseretur ad inservit

Automatic Speech Recognition (ASR)



We start with an input sequence,
like a spectrogram of audio.



The input is fed into an RNN,
for example.

h	h	h	h	h	h	h	h	h	h
e	e	e	e	e	e	e	e	e	e
l	l	l	l	l	l	l	l	l	l
o	o	o	o	o	o	o	o	o	o
ϵ									

The network gives $p_t(\alpha | X)$,
a distribution over the outputs
 $\{h, e, l, o, \epsilon\}$ for each input step.

con-
eis-
er et
nve-
mico
mo-
du-
re eu
nspa
labo-
ulip-
cidic
a. Ut
exer-
fitor
e cil-
cep-
solit
n sit

lopi
ost
nifb
cou
was
tua
whi
wate
the
an
in c
est
sev
na
nos
hi
auts
lupt
null
tem
na i

etiam
tempor
incididunt ut labore et
dolore magna aliqua. Ut enim ad minim ve-
niam, quis nostrud exercitation ullamco
labore et dolore eu ex ea commodo conser-
vatur. Duis aute irure dolor in reprehenderit
in voluptate velit esse cillum dolore eu
fugiat nulla pariatur. Excepteur sint occa-
cione
cat cupidatat non proident, sunt in culpa
qui
no
rum. Lorem ipsum sit amet, consectetur
adipiscing elit, sed do eiusmod tempor inci-
didunt ut labore et dolore magna aliqua.
ut
exer-
fitor
e cil-
cep-
solit
n sit

etiam
tempor
incididunt ut labore et
dolore magna aliqua. Ut enim ad minim ve-
niam, quis nostrud exercitation ullamco
labore et dolore eu ex ea commodo conser-
vatur. Duis aute irure dolor in reprehenderit
in voluptate velit esse cillum dolore eu
fugiat nulla pariatur. Excepteur sint occa-
cione
cat cupidatat non proident, sunt in culpa
qui
no
rum. Lorem ipsum sit amet, consectetur
adipiscing elit, sed do eiusmod tempor inci-
didunt ut labore et dolore magna aliqua.
ut
exer-
fitor
e cil-
cep-
solit
n sit

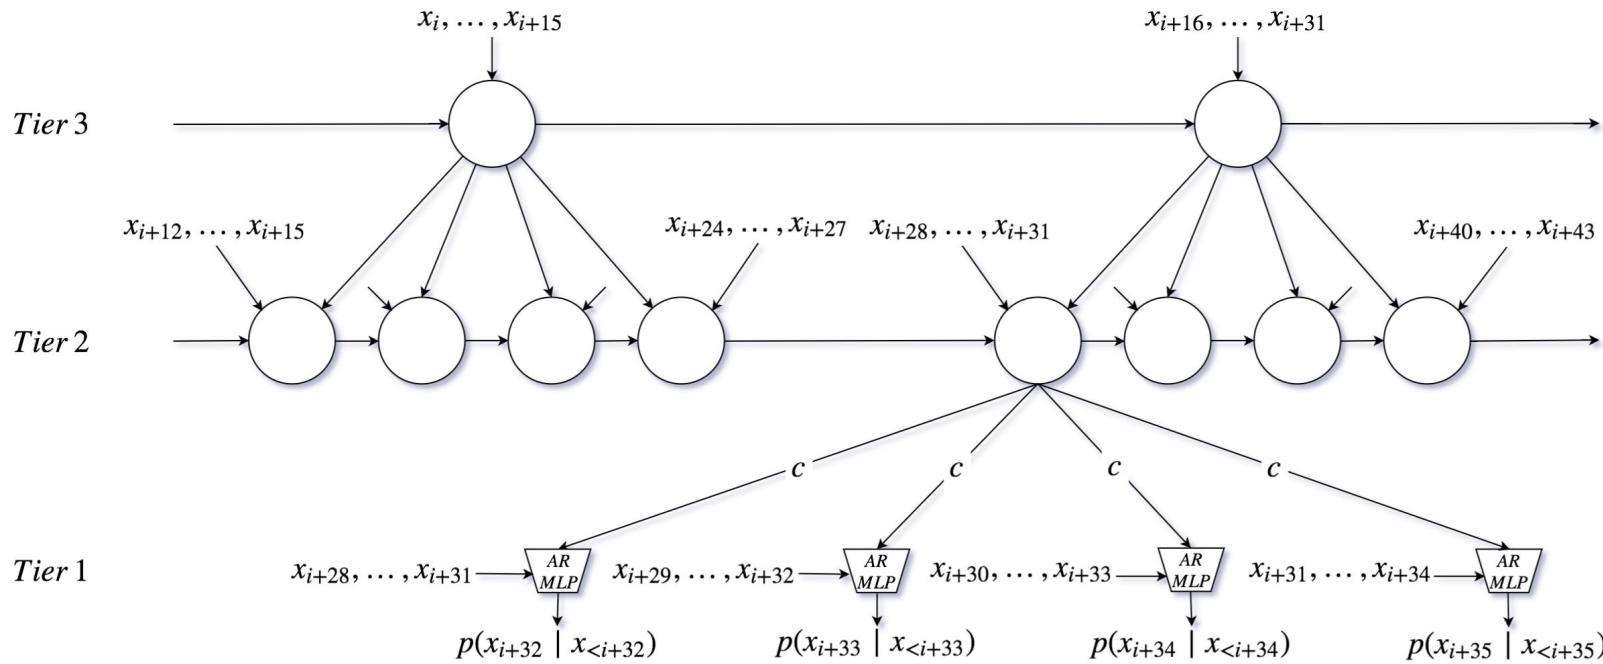
Encoder

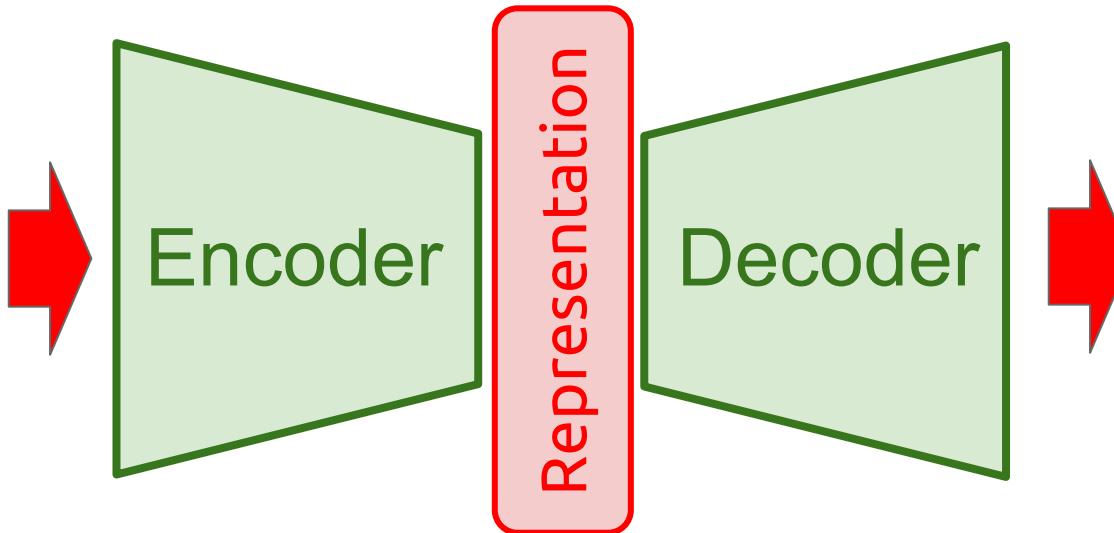
Representation

Decoder



Speech Synthesis

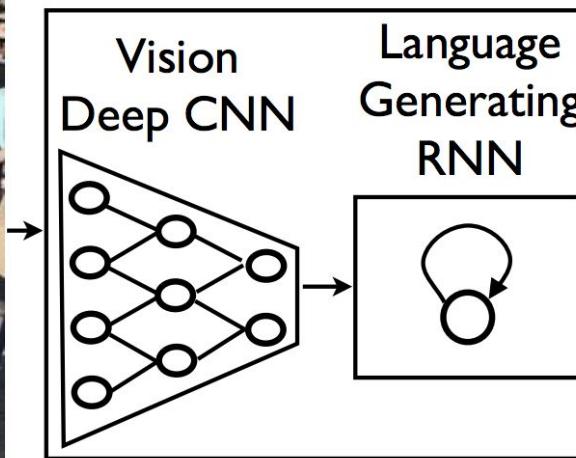




Lorem ipsum dolor

it amet, consectetur adipiscing elit,
sed do eiusmod tempor incididunt

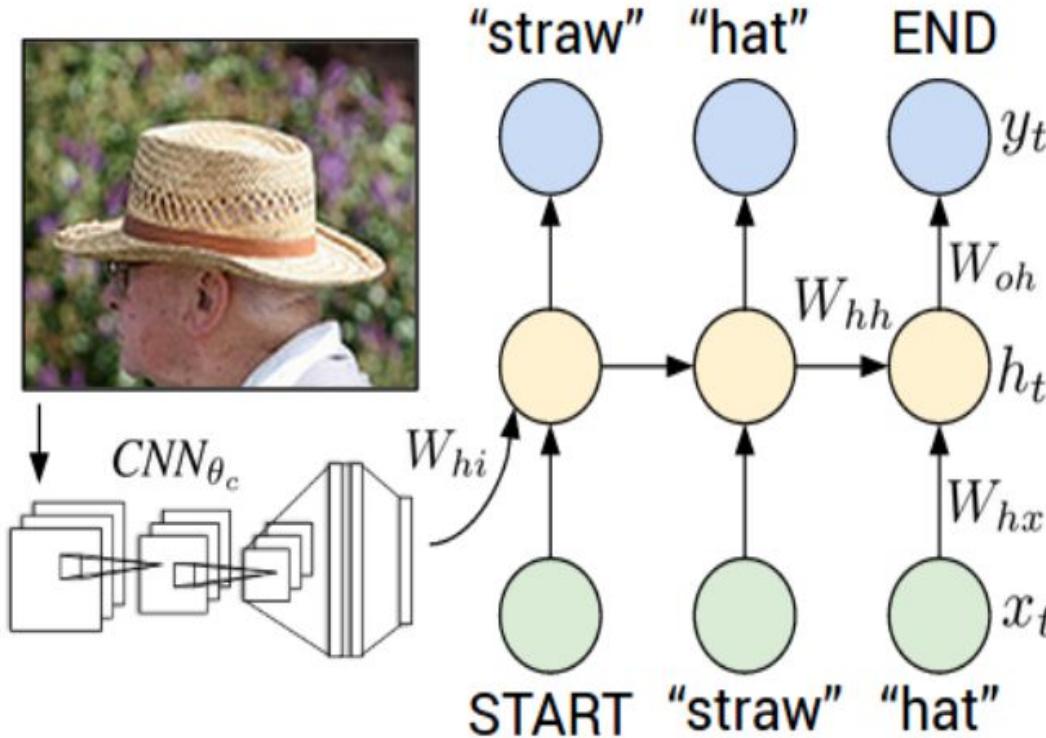
Image Captioning with RNN

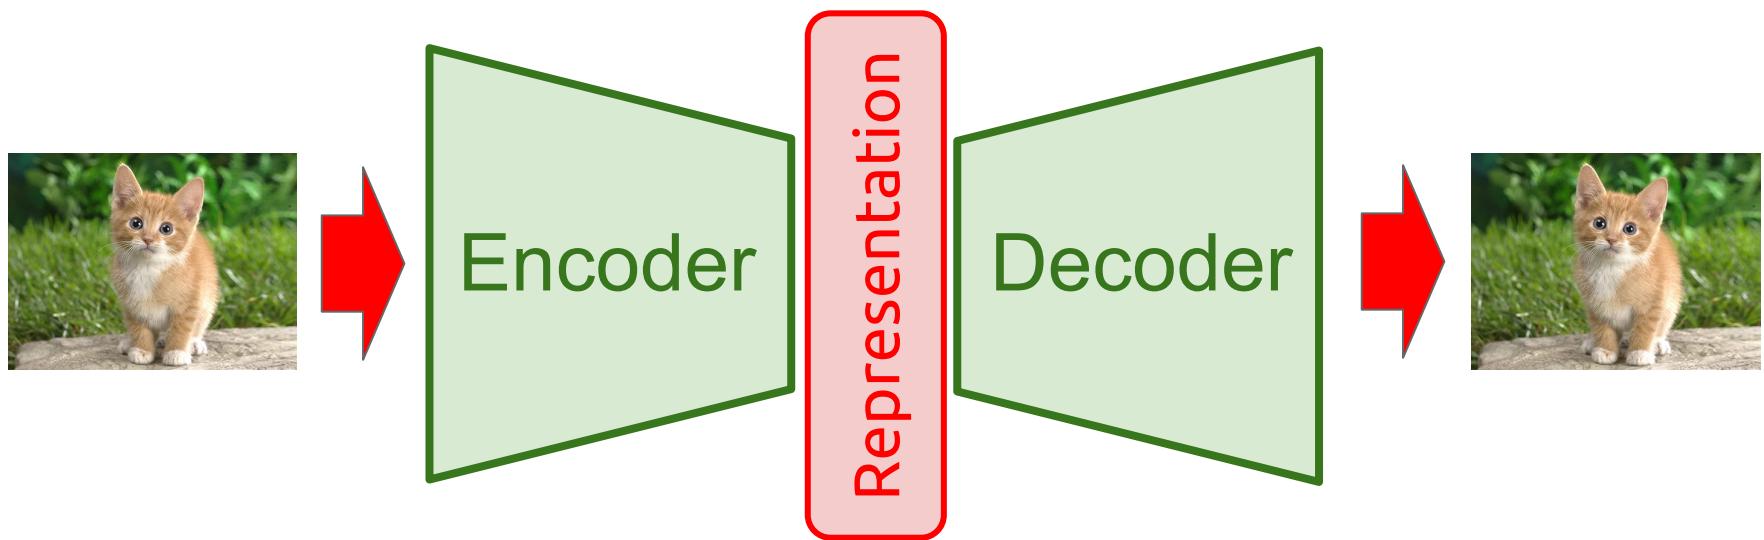


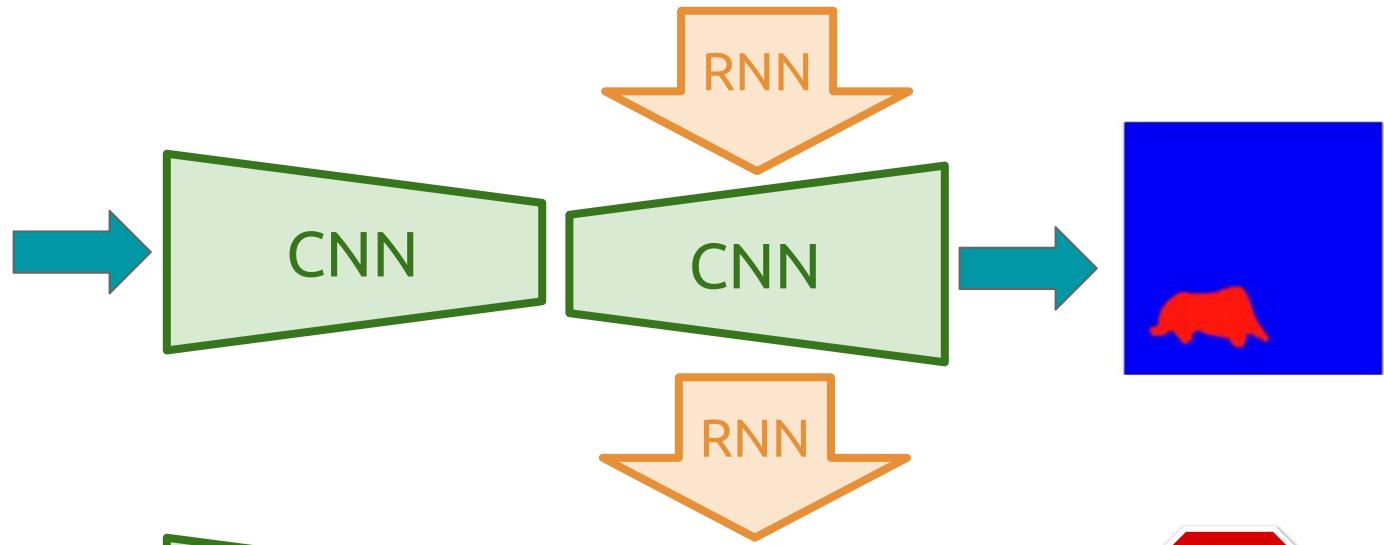
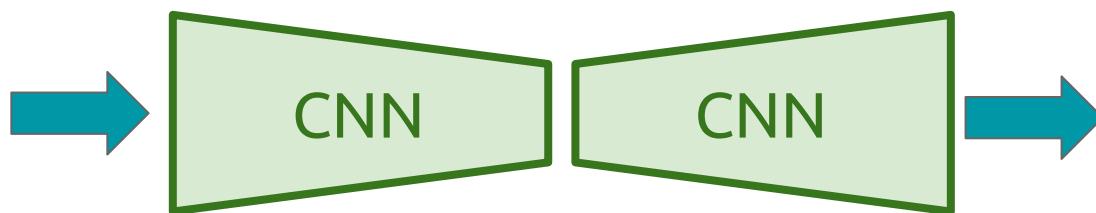
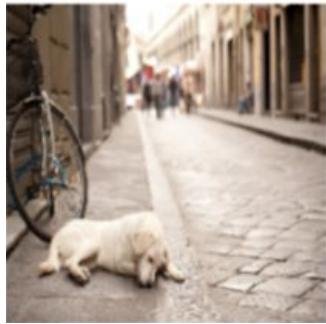
A group of people shopping at an outdoor market.

There are many vegetables at the fruit stand.

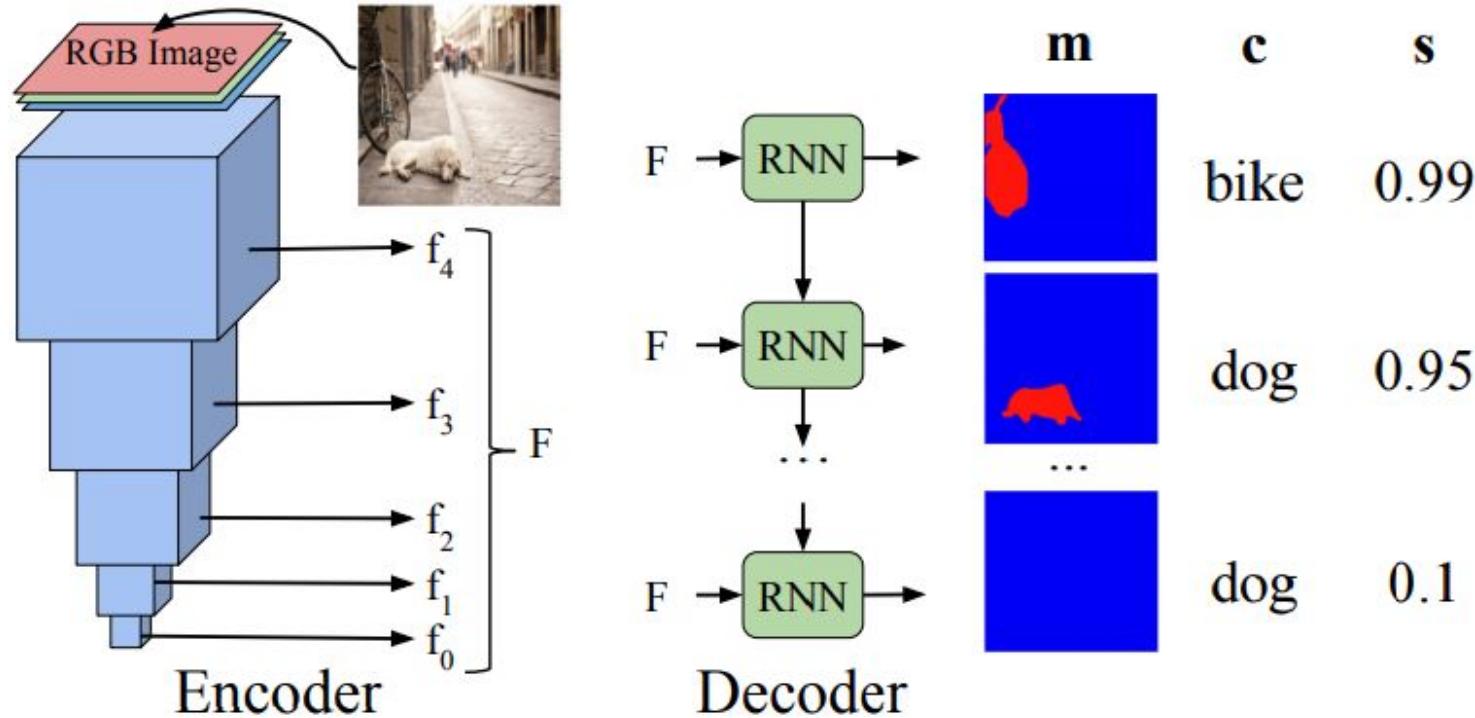
Image Captioning with RNN



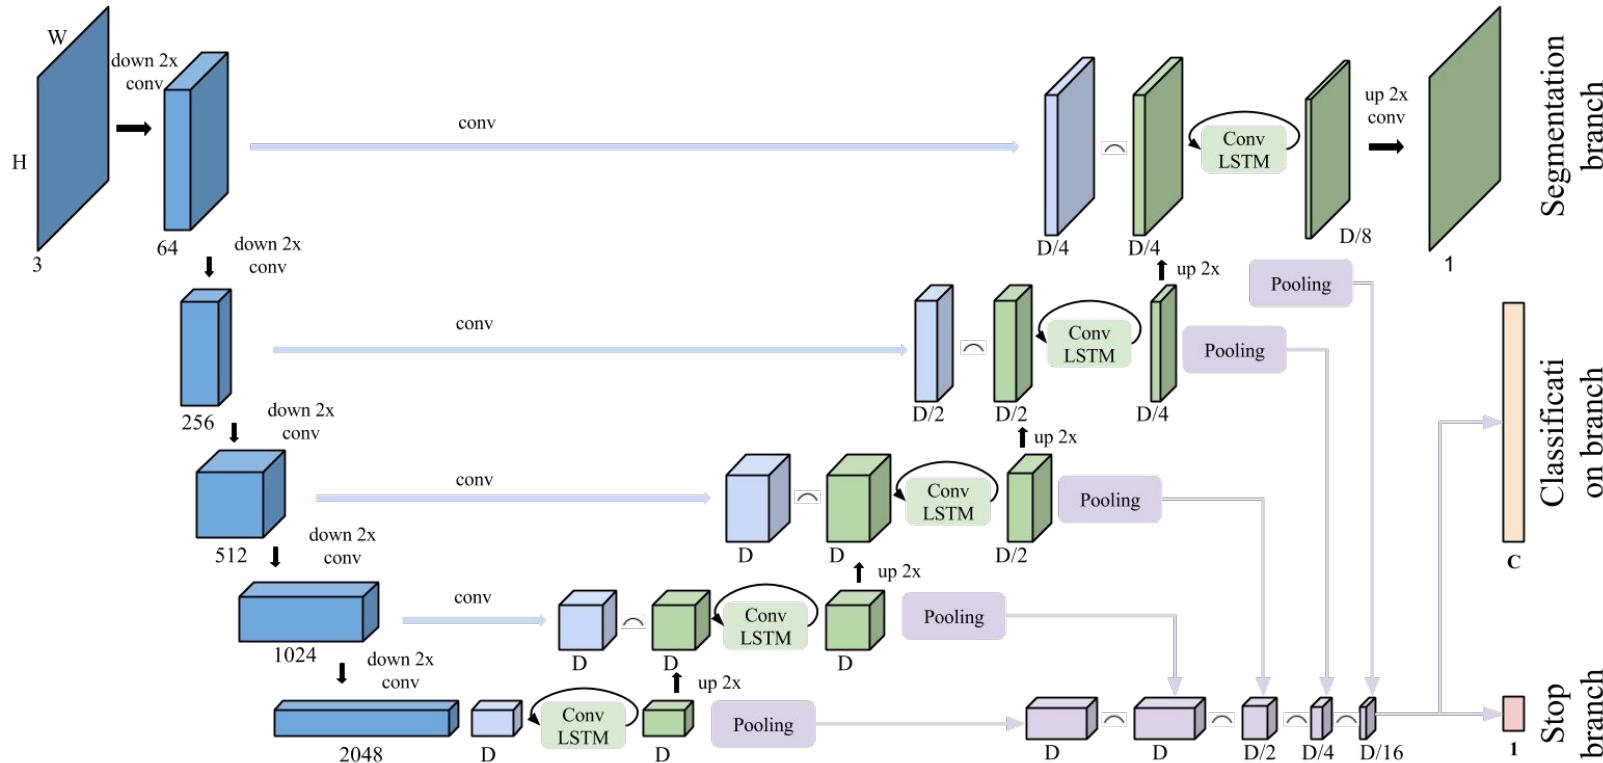




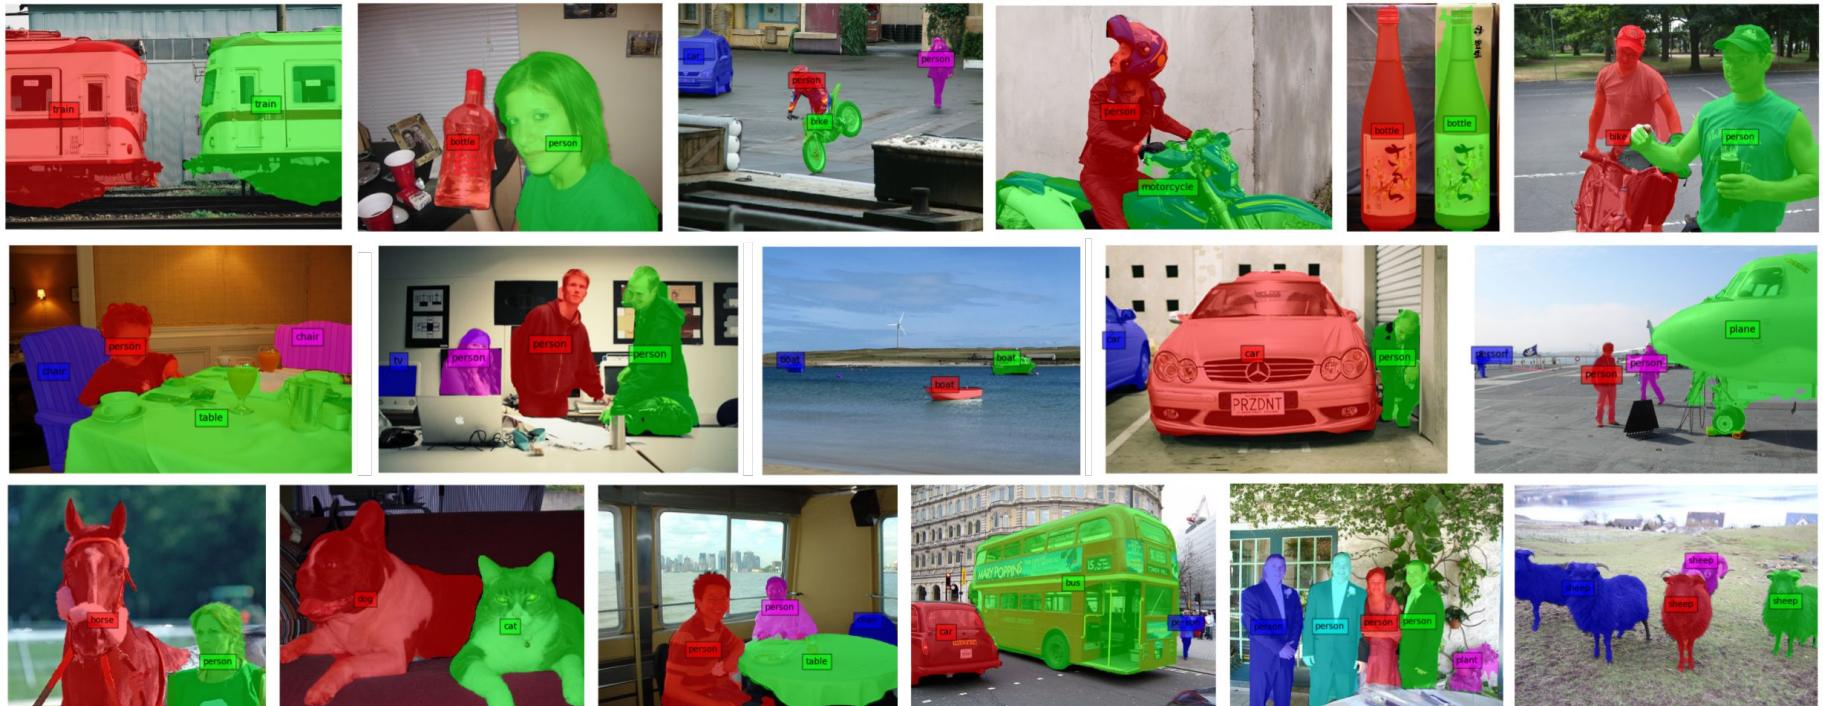
Recurrent Instance Segmentation



Recurrent Instance Segmentation



Recurrent Instance Segmentation



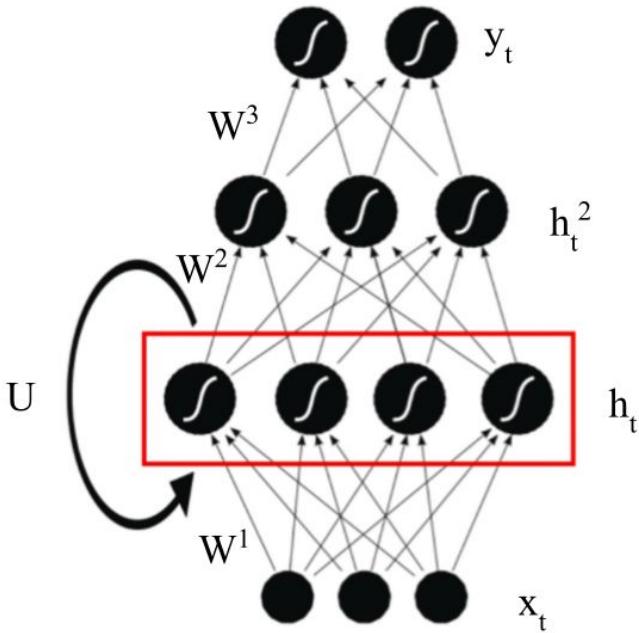
Color sequence: green red blue magenta

Outline

1. Motivation
2. Recurrent Neural Layer
3. Forward & Backward Passes Through Time
4. Gated Recurrent Neurons
 - a. LSTM
 - b. GRU
5. Advanced Architectures
6. Encoder-Decoder Arquitectures

Self-study (homework)

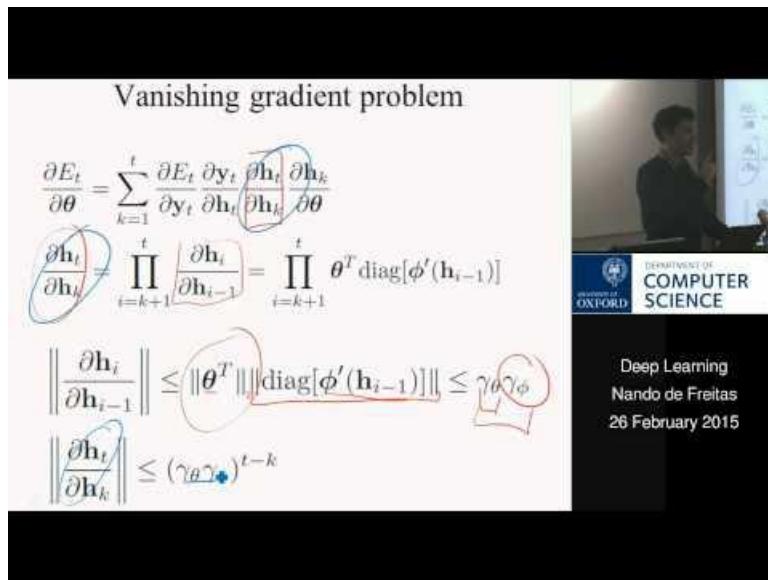
Consider a the neural network depicted in the figure, that represents the processing of a data sequence at timestep t. The layer in the box is a recurrent one, with a set of weights U.



- Formulate how is the output of the recurrent layer, h_t^1 , computed considering a $\tanh(\cdot)$ activation.
- How many parameters does this architecture contain ? Together with the depicted weight matrices W^i , also consider the biases b^i . Develop the calculations you made to reach the provided results.
- Draw the unfolded representation of this architecture for timestep $t-1$, t and $t+1$, labeling the edges with the appropriate weight matrices.

Learn more

- Chris Olah, Shan Carter, ["Attention and Augmented Recurrent Neural Networks"](#). distill.pub 2016.
- Jordi Pons (UPF): [Slides on Recurrent Neural Networks \[tweet\]](#)
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, ["Deep Learning \(Chapter 10\)"](#). MIT Press 2016.
- Alex Graves, ["Supervised Sequence Labelling with Recurrent Neural Networks"](#)



Nando de Freitas, ["Recurrent Neural Nets and LSTMs"](#) (University of Oxford 2015).



Marta Garnelo, [Deepmind x UCL 2020 \[slides\]](#)

(extra) PyTorch Lab on Google Colab

dai_2019_lab08_rnn_todo.ipynb

File Edit View Insert Runtime Tools Help Last edited on Dec 10, 2019

+ Code + Text

Recurrent Neural Networks

Lab credit

Lab created by [Santiago Pascual](#) and [Xavier Giro-i Nieto](#) for the [Postgraduate course in artificial intelligence with deep learning](#) in [UPC School \(2019\)](#).

Slides

Related slides by [Marta R. Costa-jussà](#) from [Deep Learning for Artificial Intelligence](#) (UPC TelecomBCN 2018), and slides from [Xavier Giro-i Nieto](#) from [Deep Learning for Multimedia](#) (Dublin City University 2018).

Video lectures

Related [video \(basic\)](#) & [video \(advanced\)](#) by [Santiago Pascual](#) from [Deep Learning for Speech and Language](#) (UPC TelecomBCN 2017), and [video](#) by [Marta R. Costa-jussà](#) from [Deep Learning for Artificial Intelligence](#) (UPC TelecomBCN 2019)



Santiago Pascual de la Puente
santi.pascual@upc.edu

PhD 2019

Universitat Politècnica de Catalunya
Technical University of Catalonia



DEEP LEARNING FOR ARTIFICIAL INTELLIGENCE

3rd Master Course UPC ETSETB TelecomBCN Barcelona. Autumn 2019.

Instructors

Xavier Giro-i Nieto Marta R. Costa-jussà Noé Casas Verònica Vilaplana Ramon Morros Javier Ruiz Albert Pumarola Jordi Torres

Organizers

UPC UNIVERSITAT POLITÈCNICA DE CATALUNYA BARCELONATECH telecomBCN

Supporters

Google Cloud GitHub Education

+ info: <http://bit.ly/dlai2019>

- [Lectures](#) (with Slides & Videos)
- [Labs](#)

Questions ?

Undergradese

What undergrads ask vs. what they're REALLY asking

"Is it going to be an open book exam?"

Translation: "I don't have to actually memorize anything, do I?"

"Hmm, what do you mean by that?"

Translation: "What's the answer so we can all go home."

"Are you going to have office hours today?"

Translation: "Can I do my homework in your office?"

"Can i get an extension?"

Translation: "Can you re-arrange your life around mine?"

"Is this going to be on the test?"

Translation: "Tell us what's going to be on the test."

"Is grading going to be curved?"

Translation: "Can I do a mediocre job and still get an A?"

